



# **Software Development Kit User Guide**

Contents of this document are subject to change without notice.

Cavium Proprietary and Confidential DO NOT COPY

**November 2013**

PUBLISHED BY  
Cavium, Inc.  
2315 North 1st Street  
San Jose, CA 95131  
Phone: 650-623-7000  
Fax: 650-625-9751  
Email: [sales@cavium.com](mailto:sales@cavium.com)  
Web: <http://www.cavium.com>

© 2013 by Cavium

All rights reserved. No part of this manual may be reproduced in any form, or transmitted by any means, without the written permission of Cavium.

Cavium makes no warranty about the use of its products, and reserves the right to change this document at any time, without notice. Whereas great care has been taken in the preparation of this manual, Cavium, the publisher, and the authors assume no responsibility for errors or omissions.

OCTEON is a registered trademark of Cavium.

MIPS, MIPS64, and cnMIPS are registered trademarks of MIPS Technologies. Cavium is a licensee of cnMIPS. PCI Express and PCIe are registered trademarks of PCI-SIG.

All other trademarks or service marks referred to in this manual are the property of their respective owners.

#### Revision History

Release	Date	Changes
3.1.0	November 2013	Replaced bootoct and oct-shutdown-app utilities for controlling SE applications from Linux with oct-app-ctl. Removed Valgrind sections. Added -perf3 and --perf4 command-line options to oct-profile, oct-remote-profile, and oct-top (only available in OCTEON III); updated L2 events for oct-remote-profile.

# Table of Contents

---

<b>Chapter 1 Introduction .....</b>	<b>1</b>
<b>Chapter 2 The oct-profile Tool .....</b>	<b>2</b>
2.1 Introduction .....	2
2.2 Features .....	2
2.3 Command-line Options .....	2
2.4 Getting Started .....	3
2.5 Understanding Output .....	3
<b>Chapter 3 The oct-remote-profile Tool .....</b>	<b>5</b>
3.1 Introduction .....	5
3.2 Features .....	5
3.3 Command-line Options .....	6
3.4 Getting Started .....	8
3.4.1 Running Natively on OCTEON .....	8
3.4.2 Running on OCTEON Connected over PCI .....	8
3.5 Understanding Output .....	9
<b>Chapter 4 The oct-top Tool .....</b>	<b>11</b>
4.1 Introduction .....	11
4.2 Features .....	11
4.3 Command-line Options .....	11
4.4 Getting Started .....	12
4.4.1 oct-top Commands Quick Reference .....	12
4.4.2 Running oct-top Natively on OCTEON .....	13
4.4.3 Running OCTEON Over PCI .....	14
4.5 Understanding Output .....	14
<b>Chapter 5 The otrace Tool .....</b>	<b>17</b>
5.1 Introduction .....	17
5.2 Preparing to Trace .....	18
5.3 The otrace master Commands .....	19
5.3.1 System Commands .....	19
5.3.2 Tracing Commands .....	20
5.3.3 Miscellaneous Commands .....	22
5.4 Understanding Output .....	22
5.5 Example Traces .....	24
5.5.1 Example 1 .....	24
5.5.2 Example 2 .....	25
<b>Chapter 6 The perf Tool .....</b>	<b>30</b>
6.1 Introduction .....	30
6.2 Event Mapping .....	30
6.3 Reading Additional Counters .....	32
<b>Chapter 7 The perfzilla Utility .....</b>	<b>33</b>
7.1 Introduction .....	33
7.2 Features .....	33
7.3 Getting Started .....	33
7.4 The User Interface .....	34
7.4.1 Main Menu .....	34
7.4.2 Include Cores Pane .....	35

7.4.3 perfzilla Windows .....	35
<b>Chapter 8 SE-Dump and the oct-remote-coredump Tool .....</b>	<b>39</b>
8.1 Introduction .....	39
8.2 Features .....	39
8.3 How the Tool Works .....	39
8.4 Command-line Option .....	40
8.5 Getting Started .....	40
<b>Chapter 9 The viewzilla Tool .....</b>	<b>42</b>
9.1 Introduction .....	42
9.2 Features .....	42
9.3 Getting Started .....	42
9.4 The User Interface .....	43
9.4.1 Main Menu .....	43
9.4.2 Windows and Navigation Tools .....	44
9.5 Understanding Output .....	45
9.5.1 Delay/Stall Codes .....	45
9.6 The vz-cut Utility .....	46
9.7 Tips for Running viewzilla .....	47
<b>Appendix A Controlling SE Applications from Linux .....</b>	<b>48</b>
A.1 Using the oct-app-ctl Utility .....	48
A.1.1 Starting an SE Application .....	48
A.1.2 Displaying Application Status .....	49
A.1.3 Shutting Down an Application .....	49
A.1.4 Adding Cores to an Application .....	50
A.1.5 Deleting Cores from an Application .....	50
A.2 Limitations .....	51
<b>Appendix B Examples .....</b>	<b>52</b>
B.1 oct-profile .....	52
B.1.1 Invoking oct-profile for traffic-gen .....	52
B.1.2 Profiling traffic-gen over PCI .....	52
B.2 oct-remote-profile .....	53
B.2.1 Profiling traffic-gen .....	53
B.2.2 Hardware Statistics Profiling traffic-gen .....	54
B.2.3 Profiling traffic-gen Over Subset of Cores .....	55
B.2.4 Available Events .....	55
B.2.5 Event-based Profiling for traffic-gen .....	55
B.3 oct-top .....	56
B.3.1 Frequent Routines .....	56
B.3.2 Top Information on a Core .....	56
B.3.3 PC Values .....	57
B.3.4 Frequently Called Routines and Source Code Information .....	57
B.3.5 Routines and Hexadecimal PC Values .....	58
B.3.6 Top of Selected Cores .....	58
B.3.7 OCTEON Hardware Statistics .....	59
B.3.8 Cores Running Simple Executive from Linux .....	59
B.3.9 Frequently Called Routines .....	60
B.3.10 Refresh Rate .....	60
B.4 perfzilla .....	60
B.4.1 Understanding perfzilla output .....	60
B.4.2 Source Code Window .....	60
B.4.3 Function List window .....	61
B.4.4 ASM window .....	61

## List of Figures

---

Figure 5-1	otrace master and stub for SE Application .....	17
Figure 5-2	traffic-gen Screen After Starting Trace.....	27
Figure 5-3	traffic-gen Screen After Tracing entry.....	28
Figure 5-4	traffic-gen Screen After Stopping Trace.....	29
Figure 7-1	perfzilla GUI .....	34
Figure 7-2	Include Cores Pane.....	35
Figure 7-3	File List Window.....	36
Figure 7-4	Source Code Window .....	37
Figure 7-5	Function List Window .....	37
Figure 7-6	ASM Window .....	38
Figure 9-1	viewzilla GUI .....	43
Figure B-1	perfzilla GUI .....	60

## List of Tables

---

Table 2-1	oct-profile Command-line Options.....	2
Table 2-2	oct-profile Command Quick Reference.....	3
Table 3-1	oct-remote-profile Command-line Options.....	6
Table 3-2	OCTEON II --perf1, --perf2 Events .....	6
Table 3-3	OCTEON III --perf1, --perf2, --perf3, --perf4 Events .....	7
Table 3-4	OCTEON II L2 Counter Events .....	7
Table 3-5	OCTEON III L2 Counter Events .....	7
Table 3-6	oct-remote-profile Command Quick Reference - Native OCTEON.....	8
Table 3-7	oct-remote-profile Command Quick Reference - OCTEON Connected over PCI .....	8
Table 4-1	oct-top Command-line Options.....	12
Table 4-2	oct-top Command Quick Reference.....	12
Table 4-3	L2 Statistics for TAD 0 .....	14
Table 4-4	L2 Statistic Descriptions .....	15
Table 5-1	otrace System Commands .....	19
Table 5-2	otrace Tracing Commands .....	20
Table 5-3	otrace Miscellaneous Commands .....	22
Table 6-1	Events and Register Mapping .....	30
Table 8-1	gdb session .....	41
Table 9-1	viewzilla Main Menu.....	44
Table 9-2	viewzilla Windows and Navigation Tools.....	44
Table 9-3	viewzilla: List of Delay/Stall Codes.....	46
Table A-1	Command-line Options for Booting an Application .....	49
Table A-2	Command-line Options for Shutting Down Applications.....	50
Table A-3	Command-line Options for Adding Cores .....	50
Table A-4	Command-line Options for Deleting Cores .....	51
Table B-1	oct-profile: Invoking .....	52
Table B-2	oct-profile: Profiling traffic-gen over PCI .....	52
Table B-3	oct-remote-profile: Profiling traffic-gen .....	54
Table B-4	oct-remote-profile: Hardware Statistics .....	54
Table B-5	oct-remote-profile: Profiling traffic-gen over Subset of Cores .....	55
Table B-6	oct-remote-profile: Event-based Profiling for traffic-gen.....	56
Table B-7	oct-top Frequent Routines .....	56
Table B-8	oct-top: Top Core Information .....	57
Table B-9	oct-top: PC Values in Ascending Order .....	57
Table B-10	Frequently-called Routines.....	58
Table B-11	oct-top: Routines and Hexadecimal PC Values .....	58
Table B-12	oct-top: Running utility on selected Cores .....	59
Table B-13	OCTEON Hardware Statistics .....	59
Table B-14	oct-top: Simple Executive from Linux.....	59
Table B-15	Frequently-called Routines.....	60
Table B-16	oct-top: Refresh rate.....	60
Table B-17	perfzilla: Source Code Window.....	61
Table B-18	perfzilla: Function List Window.....	61
Table B-19	perfzilla: ASM Window.....	61
Table B-20	perfzilla: ASM Window - Source Code Disassembly .....	61

# Chapter 1

## Introduction

The Cavium OCTEON Software Developer Kit (SDK) includes a set of analysis tools for Linux and Simple Executive applications. Most tools run either natively on an OCTEON target running Linux or on a remote host connected to an OCTEON board.

Native tools are provided with both the embedded root file system and Debian. The tools are usually hosted in the `/usr/bin` directory.

This guide provides instructions for using the tools, organized in the following sections:

- [The oct-profile Tool](#)
- [The oct-remote-profile Tool](#)
- [The oct-top Tool](#)
- [The otrace Tool](#)
- [The perf Tool](#)
- [The perfzilla Utility](#)
- [SE-Dump and the oct-remote-coredump Tool](#)
- [The viewzilla Tool](#)
- [Controlling SE Applications from Linux](#)

Additional examples of output are provided in [Examples](#).

# Chapter 2

## The oct-profile Tool

### 2.1 Introduction

The oct-profile utility is a profiling tool for OCTEON. It is similar to oct-remote-profile, with the key difference that oct-profile can resolve all PC values to respective routines. This facilitates the display of the results with the relevant information of routines and source code line number. An ELF file is required for profiling with oct-profile.

### 2.2 Features

The oct-profile utility includes the following features:

- Displays routine names and sample counts on respective cores.
- Displays PC (Program Count) values and associated routine information.
- Shows the source file and line number associated with PC value.
- Supports per-core view.
- Handles up to four core performance counter events.
- Handles up to four L2 (Level-2 cache) performance counters.
- Controls output refresh rate.

### 2.3 Command-line Options

The supported command-line options for oct-profile are listed in the following table:

**oct-profile** [--line] [--address] <elf>"

**Table 2-1 oct-profile Command-line Options**

Option	Description
<ELF file>	Mips ELF file from which debugging information is loaded.
--line	Display line number information.
--address	Display exact PC address.
--perf1=<Core counter>	Set up Core performance counter 1.
--perf2=<Core counter>	Set up Core performance counter 2.
--perf3=<Core counter>	Set up Core performance counter 3. <b>Note:</b> Only available in OCTEON III (CN7XXX and later).
--perf4=<Core counter>	Set up Core performance counter 4. <b>Note:</b> Only available in OCTEON III (CN7XXX and later).
--l2perf1=<L2 counter>	Set up L2 performance counter 1.



Option	Description
--l2perf2=<L2 counter>	Set up L2 performance counter 2.
--l2perf3=<L2 counter>	Set up L2 performance counter 3.
--l2perf4=<L2 counter>	Set up L2 performance counter 4.
--dwell=<seconds>	Number of seconds to refresh display.
--core_mask=<mask>	Mask that specifies the cores to sample.

## 2.4 Getting Started

The profiler can be invoked to profile applications running natively on OCTEON or applications running on an x86 development host targeting an OCTEON PCI/PCIe target board.

**Note:** The utility must always be invoked with the ELF file for the application running over PCI; it does not support profiling applications running natively on OCTEON.

The utility profiles `traffic-gen` to execute the respective commands.

**Table 2–2 oct-profile Command Quick Reference**

Command	Description
<code>oct-profile /tftpboot /traffic-gen</code>	Resolve raw PC values output to routines during runtime.
<code>oct-profile -line /tftpboot/traffic-gen</code>	Display the routines for corresponding PC values along with source code and line numbers.
<code>oct-profile -address /tftpboot/traffic-gen</code>	Display the routines for corresponding PC values along with associates PC values for that routine.
<code>oct-profile -perfl=dissue /tftpboot/traffic-gen</code>	Program the perf counter 1 for dissue event and display the results.
<code>oct-profile -dwell=2 /tftpboot/traffic-gen</code>	Wait for 2 seconds before displaying the next set of sample information.
<code>oct-profile -core_mask=0x23 /tftpboot/traffic-gen</code>	Profile cores 0,1 and 5.

**Note:**

- Post profiling tools cannot be run on the results of `oct-profile`.
- The results of `oct-profile` cannot be used as inputs for `perfzilla`.

## 2.5 Understanding Output

When `oct-profile` is invoked from an x86 host with OCTEON connected over PCI running the `traffic-gen` Simple Executive application, the output displays as shown below.

target #	oct-profile						
core#:	0	1	2	3	4	5	_sprintf_r
	1	0	0	0	0	0	_vfprintf_r
	1	0	0	0	0	0	cvmx_sysinfo_get
	1	0	0	0	0	0	fdt_offset_ptr
	2415	0	0	0	0	0	main
	2415	0	0	0	0	0	statistics_gatherer
	0	3251	3251	3251	3251	3251	packet_transmitter
	0	203	185	175	168	174	cvmx_get_cycle
	0	203	185	175	168	174	cvmx_clock_get_count
	0	91	69	62	76	87	cvmx_pow_work_request_async_nocheck
	0	20	12	6	19	16	cvmx_send_single
	0	20	12	6	19	16	cvmx_write64_uint64
	0	97	120	103	99	114	cvmx_scratch_write64
	327	0	0	0	0	0	periodic_update
	14	0	0	0	0	0	cvmx_get_cycle
	14	0	0	0	0	0	cvmx_clock_get_count
	505	0	0	0	0	0	uart_printf

The results display the sample count on respective cores for a particular routine. In the background, oct-profile retrieves the PC values and associates sample hits. It resolves all PC values to the associated routines. This makes the results more understandable than those returned by running oct-remote-profile or oct-linux-profile.

# Chapter 3

## The oct-remote-profile Tool

### 3.1 Introduction

Profiling provides a way to analyze and fix performance issues on a system. Code running on OCTEON can be profiled using oct-remote-profile.

The oct-remote-tool can run natively on OCTEON, or it can run on an x86 host targeting an OCTEON PCI/PCIe target board. The tool is invoked using a wrapper, depending on which method you are using:

- `oct-linux-profile`—the equivalent of `OCTEON_REMOTE_PROTOCOL=linux oct-remote-profile`. Use when running the tool natively on OCTEON.
- `oct-pci-profile`—the equivalent of `OCTEON_REMOTE_PROTOCOL=pci oct-remote-profile`. Use when running the tool on an x86 host targeting an OCTEON PCI/PCIe target board.

The oct-remote-profile utility supports both continuous statistical profiling and event-based profiling when run on an x86 host, while only statistical profiling is supported when running natively on OCTEON.

- **Statistical Profiler**—Profiles are collected by continuously sampling the registers on each core. The profiler collects PC values and resolves them to symbols and lines of code (LOC) in source code using post-profiling tools like `perfzilla`.
- **Event-based Profiler**—Event-based profiling takes advantage of dedicated performance counters capable of measuring events at the hardware level on each core. These counters increment once per each event and can generate interrupt on counter overflow. The oct-remote-profile utility can use these interrupts to generate samples, which are then mapped by post-profiling tools; for example, PC values to source code routine and line numbers.

### 3.2 Features

The oct-remote-profile utility includes the following features:

- Performs statistical profiling of Linux and Simple Executive applications.
- Performs interrupt-based profiling of a Simple Executive application. For similar profiling on Linux, refer to the Oprofile section (*Performance Profiling Using Oprofile*) of the documentation found under `doc/html/index.html`.
- Displays OCTEON hardware information such as memory, FPA (floating point arithmetic), and POW.
- Profiles an application on a subset of cores.
- Collects profile information specific to a core.
- Controls the profile collection rate.

- Profiles the Simple Executive application for a specific hardware event.

**Note:** Currently, OCTEON and OCTEON Plus only support profiling a single event.

### 3.3 Command-line Options

The supported command-line options for oct-remote-profile are listed in the following table:

**Table 3–1 oct-remote-profile Command-line Options**

Option	Description
--perf1	Core counter
--perf2	Core counter
--perf3	Core counter <b>Note:</b> Only available in OCTEON III (CN7XXX and later).
--perf4	Core counter <b>Note:</b> Only available in OCTEON III (CN7XXX and later).
--l2perf1	L2 counter
--l2perf2	L2 counter
--l2perf3	L2 counter
--l2perf4	L2 counter
--list-events	List available events.
--events	<event name>:<count> Enable profiling based on event and threshold count.
--fast	Sample as quickly as possible.
--samples	<samples per second> Number of samples taken per second. <b>Note:</b> High sample rates may not yield expected results.
--dwell	<seconds> Number of seconds between printing.
--perfonly	Only display performance counters.
--count	<number> Number of dwell times to sample and then exit.
--core_mask	<mask> Specify cores to mask.

The following table lists the options available for perf1 and perf2 events (OCTEON II):

**Table 3–2 OCTEON II --perf1, --perf2 Events**

clk	issue	ret	nissue
sissue	dissue	ifi	br
brmis	j	jmis	replay
iuna	trap	uunload	uustore
upload	ustore	ec	mc
cc	csrc	cfetch	cpref
ica	ii	ip	cimiss
wbuf	wdat	wbufld	wbuffl
wbuftr	badd	baddl2	bfill
ddids	idids	didna	lds

lmllds	iolds	dmllds	sts
lmsts	iosts	iobdma	dtlb
dtlbad	itlb	sync	synciob
syncw	eretmis	likmis	hazard-trap

The following table lists the options available for perf1, perf2, perf3, and perf4 events (OCTEON III):

**Table 3–3 OCTEON III --perf1, --perf2, --perf3, --perf4 Events**

none	clk	issue	ret
nissue	sissue	dissue	ifi
br	brmis	j	jmis
replay	iuna	trap	uuload
uustore	uload	ustore	ec
mc	cc	csrc	cfetch
cpref	ica	ii	ip
cimiss	wbuf	wdat	wbufld
wbuffl	wbuftr	badd	baddl2
bfill	ddids	idids	didna
lds	lmllds	iolds	dmllds
sts	lmsts	iosts	iobdma
dtlb	dtlbad	itlb	sync
synciob	syncw	dutlb	iutlb
cdmiss	eretmis	likmis	hazard-trap
fpunimp-trap	fphazard-trap	fparith-exc	fpmovcl
fpcopycl	fpcompare	fpbranch	fpmov
fpabs	fpadd	fpcvt	fpmul
fpmadd	fpdiv	fpsqrt	fpload
fpstore	fpall		

The following table lists the events available for Level 2 cache counters in OCTEON II:

**Table 3–4 OCTEON II L2 Counter Events**

none	hit	miss
no-alloc	victim	sc-fail
sc-pass	lfb-valid	lfb-wait-lfb
lfb-wait-vab	quad0-index-bus	quad0-read-bus
quad0-banks-inuse	quad0-wdat-inuse	quad1-index-bus
quad1-read-bus	quad1-banks-inuse	quad1-wdat-inuse
quad2-index-bus	quad2-read-bus	quad2-banks-inuse
quad2-wdat-inuse	quad3-index-bus	quad3-read-bus
quad3-banks-inuse	quad3-wdat-inuse	

The following table lists the events available for Level 2 cache counters in OCTEON III:

**Table 3–5 OCTEON III L2 Counter Events**

none	hit	miss	no-alloc
victim	sc-fail	sc-pass	lfb-valid
lfb-wait-lfb	lfb-wait-vab	quad0-index	quad0-read
quad0-bank	qu0-wdat	quad1-index	quad1-read
quad1-bank	quad1-wdat	quad2-index	quad2-read
quad2-bank	quad2-wdat	quad3-index	quad3-read
quad3-bank	quad3-wdat	quad4-index	quad4-read
quad4-bank	quad4-wdat	quad5-index	quad5-read
quad5-bank	quad5-wdat	quad6-index	quad6-read
quad6-bank	quad6-wdat	quad7-index	quad7-read
quad7-bank	quad7-wdat		

## 3.4 Getting Started

The profiler can be invoked to profile applications running natively on OCTEON or running on an x86 development host targeting OCTEON processors over PCI.

### 3.4.1 Running Natively on OCTEON

You can invoke `oct-linux-profile` to profile any application running on OCTEON. This includes any Linux application with an embedded root file system and Debian.

**Table 3–6 oct-remote-profile Command Quick Reference - Native OCTEON**

Command	Description
<code>oct-linux-profile</code>	Profile the code and generate PC values along with sample count for all cores.
<code>oct-linux-profile --perfonly</code>	Display only perf counters and OCTEON statistics such as TADs, DRAM memory, FPA pool memory, SSO Input/output queues.
<code>oct-linux-profile --perf1=dissue</code>	Set the perfcounter1 to dissue event and profile the application code. The output displays perfcounter1 values for the dissue event. This is a statistical profiler.
<code>oct-linux-profile --count=10</code>	Sample 10 times and exit the profiler.
<code>oct-linux-profile -core_mask=0x23</code>	Profile cores 0, 1, and 5.
<code>oct-linux-profile --dwell=2</code>	Wait two seconds before printing next set of sample data.
<code>oct-linux-profile --fast</code>	Get samples as fast as possible. Max samples are approximately 1000/sec.

### 3.4.2 Running on OCTEON Connected over PCI

Invoke `oct-pci-profile` from an x86 development host to profile applications running on OCTEON over PCI. This configuration supports both Linux and Simple Executive applications.

**Table 3–7 oct-remote-profile Command Quick Reference - OCTEON Connected over PCI**

Command	Description
<code>oct-pci-profile</code>	Profiler runs over PCI and can profile Linux or SE application.
<code>oct-pci-profile --perfonly</code>	Collect the perf counters and OCTEON statistics such as TADs, DRAM memory, FPA pool memory, SSO Input/output queues over PCI.
<code>oct-pci-profile --perf1=dissue</code>	Set perfcounter1 to dissue event and profile the application code over PCI. The output displays perfcounter1 values for the dissue event. This is a statistical profiler.
<code>oct-pci-profile --count=10</code>	Sample 10 times and exit the profiler.
<code>oct-pci-profile --core_mask=0x23</code>	Profile cores 0, 1, and 5.

Command	Description
oct-pci-profile --dwell=2	Wait two seconds before collecting next set of samples over PCI.
oct-pci-profile --fast	Get samples as fast as possible over PCI.
oct-pci-profile -list-events	List all events supported for Interrupt profiler.
oct-pci-profile -events=cimiss:100000	Invoke event-based profiler for cimiss event and the sample rate of 100000.

### 3.5 Understanding Output

When oct-pci-profile is invoked from an x86 host with OCTEON connected over PCI running the traffic-gen SE application, the output displays as shown below.

core#:	0	1	2	3	4	5
0x0000000010004f78:	7	0	0	0	0	0
0x0000000010004f8c:	1	0	0	0	0	0
0x0000000010004fe8:	267	0	0	0	0	0
0x0000000010004fec:	119	0	0	0	0	0
0x0000000010004ff0:	109	0	0	0	0	0
0x0000000010004ff8:	1	0	0	0	0	0
0x0000000010005020:	3	0	0	0	0	0
0x0000000010005028:	1	0	0	0	0	0
0x0000000010009428:	0	104	95	97	110	85
0x0000000010009430:	0	67	63	131	76	65
0x0000000010009434:	0	232	216	369	247	249
0x0000000010009630:	0	68	88	98	80	80
0x0000000010009634:	0	63	84	87	69	80
0x000000001000963c:	0	210	230	210	212	245
0x0000000010009640:	0	70	70	67	64	78
0x0000000010009644:	0	211	246	214	237	204
0x000000001000964c:	0	80	78	59	114	98
0x0000000010009734:	0	218	208	193	244	226
0x000000001000973c:	0	89	74	107	95	91
0x0000000010009740:	0	12	7	18	11	16
0x0000000010009750:	0	53	47	47	36	55
0x0000000010009754:	0	19	16	10	14	21
0x0000000010009758:	0	16	17	13	10	13
0x0000000010009778:	0	221	229	189	213	240
0x000000001000977c:	0	233	251	194	203	186
0x0000000010009784:	0	88	103	81	62	63
0x0000000010009788:	0	262	238	257	234	231
0x0000000010009ea8:	0	84	97	71	73	78
0x0000000010009eac:	0	85	82	53	77	68
0x0000000010009eb4:	0	83	68	70	69	65
0x0000000010009eb8:	0	144	122	121	134	161
0x0000000010009ebc:	0	239	222	195	267	253

The results display PC values and total samples collected for the dwell period on all available cores (for example, six cores for CN63XX).

In the sample entry, 0x0000000010009644: 0 211 246 214 237 204:

- 0x0000000010009644 is the PC value  
On OCTEON CN63XX hardware, this had:
- 0 hits on 0th core
- 211 hits on 1st core
- 246 hits on 2nd core
- 214 hits on 3rd core
- 237 hits on 4th core
- 204 hits on 5th core.

### Interrupt Profiler

For event-based profiling, the output depends on the events that would occur when profiling an application. If the sampling rate is too high, no events would occur and no samples would be collected. In such cases, oct-remote-profile displays an informative message such as the following:

```
oct-remote-profile --events=cimiss:100000000000

PCIE port 0

^C

INFO: No samples collected for 'cimiss' event with 100000000000 sample rate

INFO: Try with smaller sample-rate for same event
```



# Chapter 4

## The oct-top Tool

### 4.1 Introduction

The `oct-top` tool provides real-time information for applications running on an OCTEON target. The tool can run natively on OCTEON, or it can run on an x86 host targeting an OCTEON PCI/PCIe target board.

The `oct-top` tool generates continual reports about the state of OCTEON. The reports display the following information:

- Top CPU-consuming routines
- State of Tag and Data units (TADs)
- DRAM memory
- Floating Point Arithmetic (FPA) pool memory
- POW input/output queues

### 4.2 Features

The `oct-top` tool includes the following features:

- Prints sorted lists of frequently-called routines of applications running on an OCTEON target.
- Displays raw Program Counter (PC) values, sorted by sample count.
- Displays PC values and associated routines in sorted order.
- Shows the source file and line number associated with PC value.
- Provides per-core view.
- Handles up to four core performance counter events.
- Handles up to four L2 (Level-2 cache) performance counters.
- Controls screen refresh rate.

### 4.3 Command-line Options

The supported command-line options for `oct-top` are listed below; each option is described in the table that follows.

```
oct-top [--line] [--address] [--perfl=counter]  
[--l2perfl=L2counter] [--dwell=<seconds>] [--core_mask=mask]  
[--cpu=coreno] <elffile>
```

**Table 4–1 oct-top Command-line Options**

Option	Description
<ELF file>	Mips ELF file from which to load debugging information.
--line	Display line number information.
--address	Display exact PC address.
--perf1=<Core counter>	Set up core performance counter 1.
--perf2=<Core counter>	Set up core performance counter 2.
--perf3=<Core counter>	Set up core performance counter 3. <b>Note:</b> Only available in OCTEON III (CN7XXX and later).
--perf4=<Core counter>	Set up core performance counter 4. <b>Note:</b> Only available in OCTEON III (CN7XXX and later).
--l2perf1=<L2 counter>	Set up L2 performance counter 1. <b>Note:</b> Level-2 cache is shared by all cores and I/O components on Cavium OCTEON processors.
--l2perf2=<L2 counter>	Set up L2 performance counter 2.
--l2perf3=<L2 counter>	Set up L2 performance counter 3.
--l2perf4=<L2 counter>	Set up L2 performance counter 4.
--dwell=<seconds>	Number of seconds to refresh display.
--core_mask=<mask>	Mask that specifies the cores to sample.
--cpu=<core no>	Per-CPU top information, [0-n cores].
-r	Reverse the results of output.
-s	OCTEON Memory, Perf counter statistics.
-k	Use kernel symbols from /proc/kallsyms <b>Note:</b> This command is only supported when oct-top is run natively on OCTEON running Linux. The command can resolve the PC values to symbols (routine names), even in the absence of an ELF file. This option is not supported for Simple Executive applications.

## 4.4 Getting Started

You can invoke oct-top to run natively on OCTEON or on an x86 development host targeting the OCTEON processors over PCI. The utility can be invoked with or without an ELF file.

### 4.4.1 oct-top Commands Quick Reference

**Table 4–2 oct-top Command Quick Reference**

Command	Description
oct-top	Display PC values for all cores in ascending order
oct-top traffic-gen	Display functions which consume the most CPU time

Command	Description
<code>oct-top traffic-gen --line</code>	Display routines as well as the source code information corresponding to samples collected
<code>oct-top traffic-gen -address</code>	Display routines and associated PC values that fall under corresponding routines with sample count
<code>oct-top --dwell=2</code>	Refresh output every two seconds
<code>oct-top -core_mask=0x9 --cpu=3</code>	Display top information for core 3 where as limit the samples to coremask 0x9; that is, 0, 3 cores
<code>oct-top --cpu=4</code>	Display top information of cpu 4
<code>oct-top -r</code>	Display PC values for all cores in ascending order
<code>oct-top -s</code>	Display OCTEON statistics information such as TADs, DRAM memory, FPA pool memory, POW Input/output queues
<code>oct-top -k</code>	Use kernel symbols from <code>/proc/kallsyms</code> . Only supported when running natively on an OCTEON booted with Linux, where symbols can be resolved using <code>/proc/kallsyms</code> even in absence of an ELF file.

#### 4.4.2 Running oct-top Natively on OCTEON

Follow the instructions below to invoke oct-top from Linux running natively on OCTEON. You can run oct-top with a Linux ELF image or on a Simple Executive application booted using the oct-apt-ctl utility.

##### 1. Enable python and oct-top support.

Because oct-top has a python front end, it requires python support included in the embedded root file system and Debian. By default python support is disabled; to enable support, set the following configuration options in the `default.config` file:

```
CONFIG_python=y
CONFIG_octeon-top=y
```

You can enable the options by opening the configuration file in a text editor (for example, vi) and adding the options to the file:

```
$ vi linux/embedded_rootfs/default.config
```

or by running `make menuconfig` and selecting the options for python and octeon-top in the interface.

```
$ cd linux/embedded_rootfs
$ make menuconfig
```

##### 2. Build the kernel.

After setting the options to enable python and oct-top support, rebuild the kernel.

```
$ make [kernel|kernel-deb]
```

### 4.4.3 Running OCTEON Over PCI

All commands listed in [oct-top Commands Quick Reference](#) are valid except `oct-top -k`, which is only supported when running `oct-top` natively on an OCTEON booted with Linux.

## 4.5 Understanding Output

The following sections describe output when running `oct-top` with different options.

### `oct-top` (with no options)

When `oct-top` is invoked from an x86 host with OCTEON connected over PCI running the `traffic-gen` SE application, the output displays as shown below.

```
target# oct-top /tftpboot/traffic-gen
Function Name:                Total hits
fdt_next_tag:                  1
fdt_offset_ptr:                1
memmove:                       2
cvmx_sysinfo_get:              4
cvmx_get_cycle:                14
cvmx_clock_get_counta:         14
cvmx_send_single:              75
cvmx_write64_uint64:           75
periodic_update:               404
uart_printf:                   551
cvmx_scratch_write64:          604
cvmx_get_cycle:                1016
cvmx_clock_get_count:          1016
main:                          2613
statistics_gatherer:           2613
packet_transmitter:            17880
```

`oct-top` displays the function name and aciates samples collected for the routine during the dwell time (default is five seconds). Results are sorted in ascending order.

```
target# oct-top -s
Core Counters:
brmis 13104238[0], 12457620[1], 12186461[2], 12168803[3], 11557216[4], 275724[5],
total (excluding 0) brmis 48645824,
sync 39306[0], 39811[1], 39236[2], 38657[3], 39045[4], 26891[5],
total (excluding 0) sync 183640,
```

- Results display the default core performance counter events and counter value for dwell seconds.
- `brmis` (branch misdirects) and `sync` events are counted by default and can be changed using `-perf1` and `-perf2` options invoked with `oct-top`.
- `total` is the sum total of all performance events on available cores [0-maxcores].

### L2 Statistics

L2 refers to the Level-2 cache, which is shared by all cores and I/O components on Cavium OCTEON processors. This can be bypassed using particular CMI transactions and can also be partitioned.

CN68XX has four TADs (Tag and Data units), each containing 1 MB of L2 cache;  
CN63XX has one TAD.

**Table 4–3 L2 Statistics for TAD 0**

<code>bus_xmc(addr)</code> count:	20
<code>bus_xmd(store)</code> count:	30
<code>bus_rsc(commit)</code> count:	18
<code>bus_rsd(fill)</code> count:	20
<code>bus_ioc(IO req)</code> count:	4
<code>bus_ior(IO req)</code> count:	1

```

hit count:                191666339
miss count:               17723607
lfb-wait-lfb bus utilization: 96%
lfb-wait-vab bus utilization: 0%

```

where:

**Table 4–4 L2 Statistic Descriptions**

Statistic	Description
bus_xmc(addr) count:	L2 cache ADD bus performance counter value
bus_xmd(store) count:	L2 cache STORE bus performance counter value
bus_rsc(commit) count:	L2 cache COMMIT bus performance counter value
bus_rsd(fill) count:	L2 cache FILL bus performance counter value
bus_ioc(IO req) count:	L2 cache IOC performance counter value
bus_ior(IO req) count:	L2 cache IOR performance counter value

The following display refers to default L2 performance counter events and can be changed using `-l2perf1`, `--l2perf2`, `--l2perf3`, `--l2perf4` options:

```

hit count:                191666339
miss count:               17723607
lfb-wait-lfb bus utilization: 96%
lfb-wait-vab bus utilization: 0%

```

where:

- **LFB** = Inflight address Buffers.
- **VAB** = Victim Address Buffers

### DRAM Statistics

Includes the following statistics:

- **DRAM ops count:** This is a 64-bit Perf counter that increments each time the DDR3 databus is used.
- **dclk count:** 64-bit performance counter that increments every clock cycle.

```

DRAM ops count:    251089710
dclk count:        2667060265
utilization:       9.41%

```

### FPA Pool Statistics

Displays FPA pool statistics such as free Packet Buffers, WQE buffers, and PKO cmd buffers on OCTEON.

```

FPA pool 0 (Packet Buffers):  263 free
FPA pool 1 (WQE buffers):      860 free
FPA pool 2 (PKO cmd buffers):  941 free, PKO output queues 128

```

### POW Qos Input Queue Statistics

Queue 0:	WQE count	51	Buffers allocated	50
Queue 1:	WQE count	50	Buffers allocated	50
Queue 2:	WQE count	0	Buffers allocated	0
Queue 3:	WQE count	0	Buffers allocated	0
Queue 4:	WQE count	0	Buffers allocated	0
Queue 5:	WQE count	0	Buffers allocated	0
Queue 6:	WQE count	0	Buffers allocated	0
Queue 7:	WQE count	0	Buffers allocated	0
Total WQEs in all QOS levels:		103	(in-unit & in-memory)	
Total Free Buffers in all QOS levels:		901		

# Chapter 5

## The otrace Tool

### 5.1 Introduction

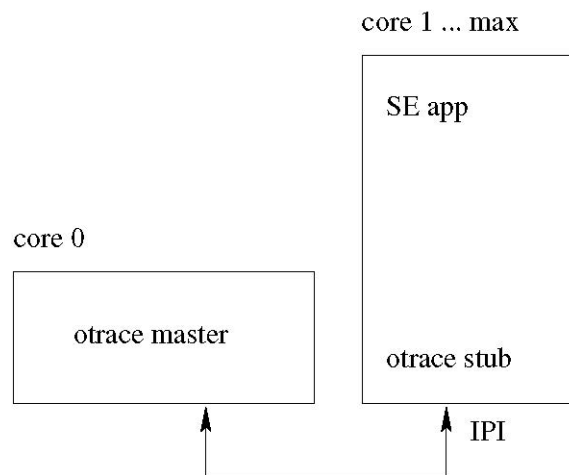
The OCTEON SE tracing tool `otrace` provides insight into a Simple Executive application's runtime behavior. The `otrace` tool has two parts:

- `stub`—Code invoked through inter-processor interrupts in the traced application.
- `otrace master`—An SE application itself that accepts user input and talks to the `stub`.

The figure below shows an example runtime setup with the `otrace master` running on one core and the application with the `stub` on the other cores.

**Note:** The `otrace master` does not have to run on core 0.

**Figure 5–1 otrace master and stub for SE Application**



The `otrace` tool was created with applications that process continuous workload in mind, for example, `passthrough`. The application and `otrace master` run concurrently during the tracing session while the user queries the application through the `otrace master` interactively. `otrace` does not explicitly support tracing applications that run to the end without waiting for user interaction (for example, `crypto`).

The following sections describe how to build an SE application to be traced and how to build a customized `otrace master` for the application. A list of `otrace master` commands is provided, as well as example tracing output.

## 5.2 Preparing to Trace

The following section provides instructions for setting up your SE application to be traced:

### 1. Prepare an SE application for `otrace`.

Embed static probes in the functions with the GCC option `-finstrument-functions` in the application's makefile:

For example:

```
OCTEON_CFLAGS_GLOBAL_ADD += -finstrument-functions \
                           -finstrument-functions-exclude-file-list= \
                               ${OCTEON_ROOT}/executive, \
                               ${OCTEON_ROOT}/target
```

In this example, the application's makefile specifies that all functions except those in the files under `executive/` and `target/` are to be instrumented.

As another example:

```
OCTEON_CFLAGS_GLOBAL_ADD += -finstrument-functions \
                           -finstrument-functions-exclude-file-list= \
                               cvmx
```

creates static probes for all functions except for those in files whose names contain `cvmx`.

**Note:** For more ways to include/exclude files and functions for instrumentation, see [gcc\(1\)](#).

A function instrumented in this manner has two probes: one at the entry and the other at the exit.

### 2. Prepare the `otrace master` for the target SE application.

The `otrace master` must know the target application's symbols and instrumentation points (static probes) to provide a more user-friendly interface. Depending on how the `otrace master` is run, this information is acquired statically or dynamically.

- **`otrace master` as an SE application**

The symbol and instrumentation information is extracted at build time and embedded in the `otrace master` binary. This is accomplished by specifying the target while building the `otrace master`.

For example:

```
TRACED_SE_APP_BIN=$(OCTEON_ROOT)/examples/passthrough/
passthrough
```

In this example, `otrace master` Makefile designates the application `passthrough` as the target application.

**Note:** In this case, both the `otrace master` and the target SE application are built with `OCTEON_TARGET=cvmx_64`.

- **`otrace master` on Linux**

When the `otrace master` runs on Linux as a userland program, there is no need or way to specify the target at build time. Instead, the `otrace master` provides



the command `ldsyms se_app_bin` to load the information at runtime from the target SE application binary `se_app_bin`.

**Note:** In this case, while the SE app is built with `OCTEON_TARGET=cvmx_64`, the `otrace master` is built with `OCTEON_TARGET=linux_64` to run on Linux.

Make sure that the target SE application is built with the symbol information.

You can now trace your application using [The otrace master Commands](#).

## 5.3 The otrace master Commands

The `otrace master` runs on one core. It provides a simple Command-Line Interface (CLI) with vi-style editing through a UART console. The `otrace` commands are divided into three groups:

- [System Commands](#)
- [Tracing Commands](#)
- [Miscellaneous Commands](#)

### 5.3.1 System Commands

The System commands include:

**Table 5–1 otrace System Commands**

Option	Description
<code>exit</code>	Exits <code>otrace master</code> .
<code>help [cmd group]</code>	Displays help information for a command or command group (tracing, system, or misc), including: <ul style="list-style-type: none"> <li>• A list of available commands with short summaries in alphabetical order when invoked without a parameter.</li> <li>• A list of commands and their summaries in a specific group with tracing, system, or miscellaneous.</li> <li>• Detailed description of a command when <code>cmd</code> supplied as the parameter.</li> </ul>
<code>ldsyms se_app_bin</code>	Loads the symbols and the probes of the program to trace.
<code>lua</code>	Switches from command-line to Lua scripting environment. See below for a description of Lua.
<code>target [core]</code>	Sets or, when run with no parameters, shows the target core.

The command “`lua`” changes the input mode. The subsequent input is directed to a Lua interpreter on a per-line basis. The user can compose and execute single-line Lua statements. Most CLI commands are available as Lua functions. To switch back to the CLI, call `cli()` in the Lua shell.

For example:

```

otrace> lua
otrace> help("system")
PP0:~CONSOLE-> 1. cli()
PP0:~CONSOLE-> Switch to cli mode.
PP0:~CONSOLE-> 2. exit()
PP0:~CONSOLE-> Exit otrace master.
PP0:~CONSOLE-> 3. help(["cmd|group"])
PP0:~CONSOLE-> Show help info for a cmd or cmd group (tracing, system, or misc).
PP0:~CONSOLE-> 4. ldsyms("se_app_bin")
PP0:~CONSOLE-> Load the symbols and the probes of the program to trace.
PP0:~CONSOLE-> 5. target([core])
PP0:~CONSOLE-> Set or, with no parameters, show the target core.
otrace> cli()
otrace> help system
PP0:~CONSOLE-> 1. exit
PP0:~CONSOLE-> Exit otrace master.
PP0:~CONSOLE-> 2. help [cmd|group]
PP0:~CONSOLE-> Show help info for a cmd or cmd group (tracing, system, or misc).
PP0:~CONSOLE-> 3. ldsyms se_app_bin
PP0:~CONSOLE-> Load the symbols and the probes of the program to trace.
PP0:~CONSOLE-> 4. lua
PP0:~CONSOLE-> Switch to lua mode.
PP0:~CONSOLE-> 5. target [core]
PP0:~CONSOLE-> Set or, with no parameters, show the target core.

```

**Note:** For more information on Lua, see R. Ierusalimsky. *Programming in Lua*. Lua.org, 2nd edition, 2006.

Many commands assume a target core, set by the command `target` with the core's number as the parameter. On the target, the core that runs the SE application under tracing cannot be the core on which the otrace master runs.

### 5.3.2 Tracing Commands

The Tracing commands include:

**Table 5–2 otrace Tracing Commands**

Option	Description
<code>fn_entry_clear func[ var1 var2 ...]</code>	Clears tracing at a function's entry on the target core.
<code>fn_entry_trace func[ var1 var2 ...]</code>	Sets tracing at a function's entry on the target core.
<code>fn_exit_clear func[ var1 var2 ...]</code>	Clears tracing at a function's exit on the target core.
<code>fn_exit_trace func[ var1 var2 ...]</code>	Sets tracing at a function's exit on the target core.
<code>shprobes [func]</code>	Shows probe info (for func) on the target core.
<code>shtrace max</code>	Displays a maximum of max log entries collected on the target core.
<code>trace [start stop]</code>	Starts or stops tracing or shows if tracing is started on the target core.

#### Probe Operations

To list the probes, run the command `shprobes` as shown in the example below.

```

otrace> shprobes
PP0:~CONSOLE-> application_shutdown
PP0:~CONSOLE->     entry off
PP0:~CONSOLE->     exit off
PP0:~CONSOLE-> application_shutdown_simple_exec
PP0:~CONSOLE->     entry off
PP0:~CONSOLE->     exit off
PP0:~CONSOLE-> main
PP0:~CONSOLE->     entry off
PP0:~CONSOLE->     exit off
PP0:~CONSOLE-> application_main_loop
PP0:~CONSOLE->     entry off
PP0:~CONSOLE->     exit off
PP0:~CONSOLE-> application_init_simple_exec
PP0:~CONSOLE->     entry off
PP0:~CONSOLE->     exit off
PP0:~CONSOLE-> afn
PP0:~CONSOLE->     entry off
PP0:~CONSOLE->     exit off
PP0:~CONSOLE-> cycle_out_of_range
PP0:~CONSOLE->     entry off
PP0:~CONSOLE->     exit off
PP0:~CONSOLE-> hello_world
PP0:~CONSOLE->     entry off
PP0:~CONSOLE->     exit off
otrace> shprobes afn
PP0:~CONSOLE-> afn
PP0:~CONSOLE->     entry off
PP0:~CONSOLE->     exit off

```

This example shows all the probe locations and those of the function `afn` upon otrace master's start. All the probes have the status `off` and are inactive.

To turn on a probe, for example, at the entry of `afn`:

```

otrace> fn_entry_trace afn
otrace> shprobes afn
PP0:~CONSOLE-> afn
PP0:~CONSOLE->     entry on
PP0:~CONSOLE->     exit off

```

When a probe is turned on, the timestamp is taken and recorded when the probe code is run.

Variables can be observed by inserting `tracers` for them in a probe. This is done using `fn_entry_trace` with the names of the variables. For example:

```

otrace> fn_entry_trace afn avar8 bvar8 cvar8
otrace> shprobes afn
PP0:~CONSOLE-> afn
PP0:~CONSOLE->     entry on
PP0:~CONSOLE->         cvar8
PP0:~CONSOLE->         bvar8
PP0:~CONSOLE->         avar8
PP0:~CONSOLE->     exit off

```

This traces three variables at the function `afn`'s entry probe. When the probe code is executed, the values of the variables are read and recorded.

These tracers can be removed from the probe using `fn_entry_clear`. The probe can be disabled using the same command with only the function name.

For example:

```

otrace> fn_entry_clear afn cvar8
otrace> shprobes afn
PP0:~CONSOLE-> afn
PP0:~CONSOLE->     entry on
PP0:~CONSOLE->         bvar8
PP0:~CONSOLE->         avar8
PP0:~CONSOLE->     exit off
otrace> fn_entry_clear afn
otrace> shprobes afn
PP0:~CONSOLE-> afn
PP0:~CONSOLE->     entry off
PP0:~CONSOLE->     exit off

```

### Tracing Operations

Once the interested events are specified, `trace start` actually kicks off tracing on the target and `shtrace max` can be executed to show the traces collected (See [Understanding Output](#)). When `trace` is executed without a parameter, it shows if tracing is enabled on the target core.

### 5.3.3 Miscellaneous Commands

The Miscellaneous commands include:

**Table 5–3 otrace Miscellaneous Commands**

Option	Description
<code>setvar var val</code>	Sets the value of a variable on the target core.
<code>shcop0 reg sel</code>	Shows the target core's COP0 register value.
<code>shvar var</code>	Shows <code>var</code> 's value on the target core.
<code>where</code>	Dumps the registers and show <code>call stack</code> on the target core.

The `where` command interrupts the target core to display its current context, including the content of the registers and the `call stack`.

For example:

```

otrace> where
r0 ($00): 0x0000000000000000    s0 ($16): 0x000000001000135c
at ($01): 0xffffffffffffbfffff    s1 ($17): 0x0000000010000000
v0 ($02): 0x0000000000000001    s2 ($18): 0x0001291de72d71f3
v1 ($03): 0x0001291de74c581b    s3 ($19): 0x0000000000000000
a0 ($04): 0x431bde82d7b634db    s4 ($20): 0x00000000100ff660
a1 ($05): 0x00000000100010c0    s5 ($21): 0xffffffffffffffff
a2 ($06): 0x000000001000135c    s6 ($22): 0x0000000000000000
a3 ($07): 0x000000000000000f    s7 ($23): 0x0000000000000000
t0 ($08): 0x0000000000000010    t8 ($24): 0x0000000000000000
t1 ($09): 0x0000000000000002    t9 ($25): 0x0000000000000017
t2 ($10): 0x0000000000000000    k0 ($26): 0x000000001001f0f8
t3 ($11): 0x0000000000000001    k1 ($27): 0xffffffff80001190
t4 ($12): 0x0000000000000000    gp ($28): 0x0000000012009180
t5 ($13): 0x0000000000000000    sp ($29): 0x00000000160fff30
t6 ($14): 0x8001180008000218    s8 ($30): 0x0000000000000000
t7 ($15): 0x00000000000d9101    ra ($31): 0x000000001000135c
    COP0_CAUSE: 0x0000000040008400
    COP0_STATUS: 0x00000000501000e1
    COP0_BADVADDR: 0x45a80f0290089683
    COP0_EPC: 0x0000000010001230
#0 0x1001049c
#1 0x1001dfe8
#2 0x1001f1c0
#3 0x10001230
#4 0x1000135c
#5 0x100006a8
#6 0x10000c48

```

**Note:** As of `otrace v1.6`, the target core does the printing.

## 5.4 Understanding Output

The `otrace stub` records runtime events and passes them to the `otrace master` as logs through a per-core ring buffer. As in a single producer-consumer model, the `stub` continues to append to the buffer until it is full, at which point it records a buffer full event. The `stub` continues to update the buffer full event until the `otrace master` consumes logs and makes room in the buffer.

As of otrace v1.6, each log entry obtained by the shtrace command contains the following items:

- **event type**—Identifies four types of events:
  - entry event generated upon entering a probed function.
  - exit event generated upon exiting a probed function.
  - variable event generated for a variable tracer in a probe.
  - buffer full event occurs when the buffer between the otrace master and the stub is full.
- **time stamp**—The cycle count upon the event.
- **type-specific content**—For function entry and exit events, contains the instrumented function's address and the address where it is called (for example, its call site). For buffer full events, contains the timestamp for the most recent attempt to record a log and the number of failed attempts so far due to buffer full.

Conceptually, a variable tracer resides in a probe. Therefore, the timestamp is only taken on the probe.

The output of an example trace is displayed below:

```
otrace> fn_entry_trace afn avar8 bvar8 cvar8
otrace> shprobes afn
PP0:~CONSOLE-> afn
PP0:~CONSOLE->      entry on
PP0:~CONSOLE->      cvar8
PP0:~CONSOLE->      bvar8
PP0:~CONSOLE->      avar8
PP0:~CONSOLE->      exit off
otrace> trace start
otrace> shtrace 10
fn <afn(0)> caller <application_main_loop(6740)> event <fn entry> [12858s:816432426(cycles)]
var addr <0x12005a88> value <0x0>
var addr <0x12005a78> value <0x0>
var addr <0x12005a68> value <0x0>
fn <afn(0)> caller <application_main_loop(6740)> event <fn entry> [12858s:822511218(cycles)]
var addr <0x12005a88> value <0x0>
var addr <0x12005a78> value <0x0>
var addr <0x12005a68> value <0x0>
fn <afn(0)> caller <application_main_loop(6740)> event <fn entry> [12858s:828588585(cycles)]
var addr <0x12005a88> value <0x0>
otrace> shtrace 1022
...
fn <afn(0)> caller <application_main_loop(6740)> event <fn entry> [12859s:186646953(cycles)]
var addr <0x12005a88> value <0x0>
event <buf full> [12862s:935655489(cycles)] starting [12859s:186647128(cycles)] #lost events=6926
fn <afn(0)> caller <application_main_loop(6740)> event <fn entry> [12862s:935655521(cycles)]
var addr <0x12005a88> value <0x0>
var addr <0x12005a78> value <0x0>
var addr <0x12005a68> value <0x0>
fn <afn(0)> caller <application_main_loop(6740)> event <fn entry> [12862s:938358057(cycles)]
var addr <0x12005a88> value <0x0>
var addr <0x12005a78> value <0x0>
var addr <0x12005a68> value <0x0>
fn <afn(0)> caller <application_main_loop(6740)> event <fn entry> [12862s:941060649(cycles)]
```

In these results, for each function entry event:

- afn(0) is the instrumented function (afn is the function name and (0) is the offset).
- application\_main\_loop(6740) is the caller.
- secs:cycles is the format for the timestamp.

Under each function entry event, the three traced variables' addresses and values are displayed. For the buffer full event, the two timestamps represent the last and first time where buffer full occurred; the #lost events value indicates 6926 events were lost during buffer full.

## 5.5 Example Traces

The following sections provide step-by-step examples of using otrace.

### 5.5.1 Example 1

1. Define the following in **passthrough.c** and build the application.

```
static uint8_t  foo8  __attribute__((used));
static uint16_t foo16 __attribute__((used));
static uint32_t foo32 __attribute__((used));
static uint64_t foo64 __attribute__((used));
static uint8_t  bar8  __attribute__((used));
static uint16_t bar16 __attribute__((used));
static uint32_t bar32 __attribute__((used));
static uint64_t bar64 __attribute__((used));
static void hello_world(void) __attribute__((used));      static void
hello_world(void)
{
    cvmx_wait_usec(2000);
}
```

2. Launch otrace on core 0 and the application on other cores (on an OCTEON CN63xx).

```
Octeon ep6300c# tftp 0 passthrough; bootoct 0 coremask=3e
Using octmgt0 device
TFTP from server 13.14.15.1; our IP address is 13.14.15.2
Filename 'passthrough'.
Load address: 0x20000000
Loading: #####
done
Bytes transferred = 4469099 (44316b hex)
Bootloader: Done loading app on coremask: 0x3e
Octeon ep6300c# tftp 0 otrace; bootoct 0 coremask=1
Using octmgt0 device
TFTP from server 13.14.15.1; our IP address is 13.14.15.2
Filename 'otrace'.
Load address: 0x20000000
Loading:
#####
done
Bytes transferred = 5507607 (540a17 hex)
Bootloader: Done loading app on coremask: 0x1
eP0P:P1~:CO~NCOSNOSLOEL-E> -U>s iUsnign gde vdiecviec etr teer
e
PP1:~CONSOLE-> Version: Cavium Inc. OCTEON SDK version 3.0.0, build 469
PP1:~CONSOLE-> Enabling CVMX_IPD_CTL_STATUS[NO_WPTR]
PP1:~CONSOLE-> Interface 0 has 4 ports (SGMII)
PP1:~CONSOLE-> Interface 3 has 4 ports (LOOP)
```

3. Complete the following steps in otrace.

- a. Check the probe status in `hello_world()`.

```
otrace> shprobes hello_world
PP0:~CONSOLE-> hello_world
PP0:~CONSOLE->      entry off
PP0:~CONSOLE->      exit
off
```

- b. Activate the entry probe and exit probe with variables.

```
otrace> fn_entry_trace hello_world
otrace> fn_exit_trace hello_world foo16
otrace> fn_exit_trace hello_world bar64
otrace> shprobes hello_world
PP0:~CONSOLE-> hello_world
PP0:~CONSOLE->      entry on
PP0:~CONSOLE->      exit on
PP0:~CONSOLE->      bar64
PP0:~CONSOLE->      foo16
```

No trace collected at this point as `hello_world()` is not called anywhere.

```
otrace> trace start
otrace> shtrace 10
```

- c. Call `hello_world()` on the target core from `otrace` master.

```
otrace> crosscall hello_world
```

- d. Review results from having probes executed once.

```
otrace> shtrace 10
fn <hello_world(0)> caller <otc_call(dc)> event <fn entry>
[326s:26011579(cycles)]
fn <hello_world(0)> caller <otc_call(dc)> event <fn exit>
[326s:28512046(cycles)]
var addr <0x12005a34> value <0x0>
var addr <0x12005a18> value <0x0>
```

- e. Set the value of `bar64`.

```
otrace> setvar bar64 1234
otrace> shvar bar64
PP0:~CONSOLE-> 0x4d2
```

- f. Call `hello_world()` again.

```
otrace> crosscall hello_world
otrace> shtrace 20
fn <hello_world(0)> caller <otc_call(dc)> event <fn entry>
[609s:40795486(cycles)]
fn <hello_world(0)> caller <otc_call(dc)> event <fn exit>
[609s:43295641(cycles)]
var addr <0x12005a34> value <0x0>
var addr <0x12005a18> value <0x4d2>
```

## 5.5.2 Example 2

1. Build `passthrough` and `otrace-linux_64`.

Build `passthrough` (`cvmx_64`) with `-finstrument-functions`. This instruments virtually everything. Build the `otrace` master with `OCTEON_TARGET=linux_64`.

2. Set up.

```
cn52xx          <-----> cn63xx
(traffic-gen)    4 GBE ports back-to-back    (passthrough on 3 cores
                                                and
                                                Linux + otrace on core 0)
```

### 3. Boot the target.

```
tftp 0 passthrough; bootoct 0 coremask=0xe; tftp 0 vmlinux.64; bootoctlinux 0
coremask=1
```

### 4. Start otrace after shipping otrace-linux\_64 and passthrough to Linux,

```
~ # chmod 700 otrace-linux_64
~ # ./otrace-linux_64
CVMX_SHARED: 0x1202f0000-0x120320000
Active coremask = 0x1
```

### 5. Import passthrough's symbols into otrace master.

```
otrace> ldsyms passthrough
```

### 6. Review the probes of cvmx\_wqe\_get\_port( ).

```
otrace> shprobes cvmx_wqe_get_port@0x10001ba0
cvmx_wqe_get_port@0x10001ba0
    entry off
    exit off
otrace> shprobes cvmx_wqe_get_port@0x10013cd8
cvmx_wqe_get_port@0x10013cd8
    entry off
    exit off
```

The 0x10001ba0 and 0x10013cd8 are addresses of two instances of the inline function `cvmx_wqe_get_port( )`. One can obtain these using `nm`; for example:

```
% mipsisa64-octeon-elf-nm passthrough|grep cvmx_wqe_get_port
10001ba0 t cvmx_wqe_get_port
10013cd8 t cvmx_wqe_get_port
```

You can also obtain these using the `otrace shprobes` command; for example:

```
otrace> shprobes wqe_get_port
PP0:~CONSOLE-> cvmx_wqe_get_port@0x10013cd8
PP0:~CONSOLE->    entry off
PP0:~CONSOLE->    exit off
PP0:~CONSOLE-> cvmx_wqe_get_port@0x10001ba0
PP0:~CONSOLE->    entry off
PP0:~CONSOLE->    exit off
```

### 7. Activate the probes at the entry and exit of both instances.

```
otrace> fn_entry_trace cvmx_wqe_get_port@0x10001ba0
otrace> fn_entry_trace cvmx_wqe_get_port@0x10013cd8
otrace> fn_exit_trace cvmx_wqe_get_port@0x10001ba0
otrace> fn_exit_trace cvmx_wqe_get_port@0x10013cd8
otrace> shprobes cvmx_wqe_get_port@0x10013cd8
cvmx_wqe_get_port@0x10013cd8
    entry on
    exit on
otrace> shprobes cvmx_wqe_get_port@0x10001ba0
cvmx_wqe_get_port@0x10001ba0
    entry on
    exit on
```



## 8. Start tracing.

```
otrace> trace start
otrace> shtrace 10
otrace>
```

Nothing returns immediately, so start pumping packets in traffic-gen:

```
otrace> shtrace 10
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> \
  event <fn entry> [521s:1020886065(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> \
  event <fn exit> [521s:1020886657(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> \
  event <fn entry> [521s:1020897379(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> \
  event <fn exit> [521s:1020897733(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> \
  event <fn entry> [521s:1020907514(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> \
  event <fn exit> [521s:1020907883(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> \
  event <fn entry> [521s:1020917533(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> \
  event <fn exit> [521s:1020917878(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> \
  event <fn entry> [521s:1020927528(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> \
  event <fn exit> [521s:1020927897(cycles)]
```

**Figure 5-2 traffic-gen Screen After Starting Trace**

0:00:23	81181	*Port 0	*Port 1	*Port 2	*Port 3	Totals
23:RX packets		186938	187204	187204	187204	748550
24:RX octets		92160370	92291500	92291519	92291548	369034937
26:RX Mbps		767	768	768	768	3071
27:tx.size		489	489	489	489	
28:tx.type		IPv4+UDP	IPv4+UDP	IPv4+UDP	IPv4+UDP	
29:tx.payload		abc	abc	abc	abc	
30:TX packets		243632	243632	243633	243632	974529
31:TX octets		120110576	120110576	120110580	120110576	480442308
33:TX Mbps		1000	1000	1000	1000	4000
35:Total RX packets		4263290	4264191	4268868	4267815	17064164
38:Total TX packets		5520354	5520356	5520358	5520359	22081427
47:dest.mac		fb71046a5	fb71046a6	fb71046a7	fb71046a4	
51:src.mac		fb71046a4	fb71046a5	fb71046a6	fb71046a7	
55:dest.ip		10.1.0.99	10.2.0.99	10.3.0.99	10.0.0.99	
59:src.ip		10.0.0.99	10.1.0.99	10.2.0.99	10.3.0.99	
73:bridge		off	off	off	off	
77:validate packets		off	off	off	off	

## 9. Trace only the entry.

```

otrace> trace stop
otrace> fn_exit_clear cvmx_wqe_get_port@0x10001ba0
otrace> shprobes cvmx_wqe_get_port@0x10001ba0
cvmx_wqe_get_port@0x10001ba0
    entry on
    exit off
otrace> trace start
otrace> shtrace 10
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> event \
    <fn entry> [699s:856353998(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> event \
    <fn entry> [699s:856363993(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> event \
    <fn entry> [699s:856373962(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> event \
    <fn entry> [699s:856383922(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> event \
    <fn entry> [699s:856393880(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> event \
    <fn entry> [699s:856403863(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> event \
    <fn entry> [699s:856413820(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> event \
    <fn entry> [699s:856423778(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> event \
    <fn entry> [699s:856433736(cycles)]
fn <cvmx_wqe_get_port@0x10001ba0(0)> caller <application_main_loop(1dc)> event \
    <fn entry> [699s:856443718(cycles)]

```

**Figure 5-3 traffic-gen Screen After Tracing entry**

0:03:52	81390	*Port 0	*Port 1	*Port 2	*Port 3	Totals
23:RX packets		187514	187988	187747	187988	751237
24:RX octets		92443994	92678244	92559303	92678164	370359705
26:RX Mbps		770	772	771	772	3085
27:tx.size		489	489	489	489	
28:tx.type		IPv4+UDP	IPv4+UDP	IPv4+UDP	IPv4+UDP	
29:tx.payload		abc	abc	abc	abc	
30:TX packets		243632	243632	243632	243632	974528
31:TX octets		120110576	120110576	120110576	120110576	480442304
33:TX Mbps		1000	1000	1000	1000	4000
35:Total RX packets		48150614	48169472	48182670	48182220	192684976
38:Total TX packets		56439474	56439476	56439478	56439479	225757907
47:dest.mac		fb71046a5	fb71046a6	fb71046a7	fb71046a4	
51:src.mac		fb71046a4	fb71046a5	fb71046a6	fb71046a7	
55:dest.ip		10.1.0.99	10.2.0.99	10.3.0.99	10.0.0.99	
59:src.ip		10.0.0.99	10.1.0.99	10.2.0.99	10.3.0.99	
73:bridge		off	off	off	off	
77:validate packets		off	off	off	off	

## 10. Stop the trace.

```

otrace> trace stop
otrace>

```

Figure 5–4 traffic-gen Screen After Stopping Trace

0:04:29	81427	*Port 0	*Port 1	*Port 2	*Port 3	Totals
23:RX packets		243632	243633	243633	243643	974541
24:RX octets		120110736	120110733	120110741	120116127	480448337
26:RX Mbps		1000	1000	1000	1000	4000
27:tx.size		489	489	489	489	
28:tx.type		IPv4+UDP	IPv4+UDP	IPv4+UDP	IPv4+UDP	
29:tx.payload		abc	abc	abc	abc	
30:TX packets		243633	243632	243632	243633	974530
31:TX octets		120111069	120110576	120110576	120111069	480443290
33:TX Mbps		1000	1000	1000	1000	4000
35:Total RX packets		56336551	56356045	56370284	56371612	225434492
38:Total TX packets		65453865	65453867	65453869	65453871	261815472
47:dest.mac		fb71046a5	fb71046a6	fb71046a7	fb71046a4	
51:src.mac		fb71046a4	fb71046a5	fb71046a6	fb71046a7	
55:dest.ip		10.1.0.99	10.2.0.99	10.3.0.99	10.0.0.99	
59:src.ip		10.0.0.99	10.1.0.99	10.2.0.99	10.3.0.99	
73:bridge		off	off	off	off	
77:validate packets		off	off	off	off	

# Chapter 6

## The perf Tool

### 6.1 Introduction

Perf is a profiling tool for Linux 2.6 and higher-based systems based on the `perf_events` interface. The tool provides abstractions for hardware-specific capabilities to measure performance through a simple command-line interface.

This section provides specific information about the mapping between perf hardware events and OCTEON Performance Control Registers. For a complete description of how to use, refer to the tutorial on the perf wiki (<https://perf.wiki.kernel.org/index.php/Tutorial>).

### 6.2 Event Mapping

Perf maps hardware events to OCTEON Performance Control Registers as described in the following table:

**Table 6–1 Events and Register Mapping**

Perf Event	OCTEON Performance Control Register	Description (supported events from OCTEON HRM)
<code>cpu-cycles</code> OR <code>cycles</code>	[PERF_CNT_CLK] (reported in HZ)	Conditionally clocked cycles (as opposed to <code>count/cvm_count</code> , which counts even with no clocks)
<code>stalled-cycles-frontend</code> OR <code>idle-cycles-frontend</code>	N/A	
<code>stalled-cycles-backend</code> OR <code>idle-cycles-backend</code>	N/A	
<code>instructions</code>	[PERF_CNT_RET]	Instructions retired
<code>cache-references</code>	[PERF_CNT_LDS]	Number of load issues
<code>cache-misses</code>	[PERF_CNT_DMLDS]	Number of loads that were not prefetched and missed in the cache
<code>branch-instructions</code> OR <code>branches</code>	[PERF_CNT_BR]	Branches retired, excluding branch likelies
<code>branch-misses</code>	[PERF_CNT_BRMIS]	Branch mispredicts, excluding branch likelies
<code>bus-cycles</code>	[PERF_CNT_BADD]	Number of address bus cycles used (may need to set <code>CvmCtl[DISCE]</code> for accurate counts)
<code>ref-cycles</code>	N/A	

Perf Event	OCTEON Performance Control Register	Description (supported events from OCTEON HRM)
L1-dcache-loads	[PERF_CNT_LDS]	Number of load issues
L1-dcache-load-misses	[PERF_CNT_DMLDS]	Number of loads that were not prefetched and missed in the cache
L1-dcache-stores	[PERF_CNT_STS]	Number of store issues
L1-dcache-store-misses	N/A	
L1-dcache-prefetches	N/A	
L1-dcache-prefetch-misses	N/A	
L1-icache-loads	[PERF_CNT_CFETCH]	Icache committed fetches (demand+prefetch)
L1-icache-load-misses		
L1-icache-prefetches	[PERF_CNT_CPREF]	Icache committed prefetches
L1-icache-prefetch-misses	N/A	
LLC-loads	N/A	
LLC-load-misses	N/A	
LLC-stores	N/A	
LLC-store-misses	N/A	
LLC-prefetches	N/A	
LLC-prefetch-misses	N/A	
dTLB-loads	[PERF_CNT_DTLB]	Number of dstream TLB refill, invalid, or modified exceptions
dTLB-load-misses	[PERF_CNT_DTLB]	Number of dstream TLB refill, invalid, or modified exceptions
dTLB-stores	[PERF_CNT_DTLB]	Number of dstream TLB refill, invalid, or modified exceptions
dTLB-store-misses	[PERF_CNT_DTLB]	Number of dstream TLB refill, invalid, or modified exceptions
dTLB-prefetches	[PERF_CNT_DTLB]	Number of dstream TLB refill, invalid, or modified exceptions
dTLB-prefetch-misses	[PERF_CNT_DTLB]	Number of dstream TLB refill, invalid, or modified exceptions
iTLB-loads	[PERF_CNT_ITLB]	Number of dstream TLB refill, invalid, or modified exceptions
iTLB-load-misses	[PERF_CNT_ITLB]	Number of dstream TLB refill, invalid, or modified exceptions

Perf Event	OCTEON Performance Control Register	Description (supported events from OCTEON HRM)
branch-loads	[PERF_CNT_J]	Jumps retired; does not include J/JAL
branch-load-misses	[PERF_CNT_JMIS]	Jumps mispredicted; does not include J/JAL
node-loads	N/A	
node-load-misses	N/A	
node-stores	N/A	
node-store-misses	N/A	
node-prefetches	N/A	
node-prefetch-misses	N/A	

## 6.3 Reading Additional Counters

Some OCTEON counters are not mapped to perf. These can be read from `/proc/octeon_perf`.

**Important!** `/proc/octeon_perf` and `perf` step over each other when reading counters; consequently, they should not be used simultaneously.

# Chapter 7

## The perfzilla Utility

### 7.1 Introduction

The perfzilla utility is a post-profiling graphical tool that analyzes the results from `oct-remote-profile` or `oct-linux-profile`. It derives meaningful information from the profile logs collected over a longer period of time using the profiling tools.

### 7.2 Features

The perfzilla utility includes the following features:

- Provides a Graphical User Interface (GUI).
- Displays the sorted lists of files and functions in descending order based on the Hit Count in the **File List** window.
- Displays the sample percentage and sample count against the file names in the **File List** window.
- Displays the source code of the corresponding file in the **Source Code** window.
- When a file is double clicked, displays the first hit count in the **Source Code** window; displays subsequent hits on every double click.
- Displays view of **Next Hit** and **Previous Hit** by selecting the corresponding action from **View** menu.
- Supports per core/set of cores sample views.
- Supports copy, paste, and search options.
- Supports mouse wheel option for all text widgets and scroll bars in the window.

### 7.3 Getting Started

The perfzilla utility requires the following two files as input:

- `profile.log`—Contains the captured output of `oct-pci-profile` or a Viewzilla trace.
- `ELF file`—Contains Dwarf2 debugging information.

To start perfzilla, run the following commands:

```
$ sudo oct-remote-profile > profile.log
$ perfzilla profile.log elf_file
```

This prompts perfzilla to complete the following steps:

- Load the ELF symbols.
- Load the line information of supplied `elf_file`.



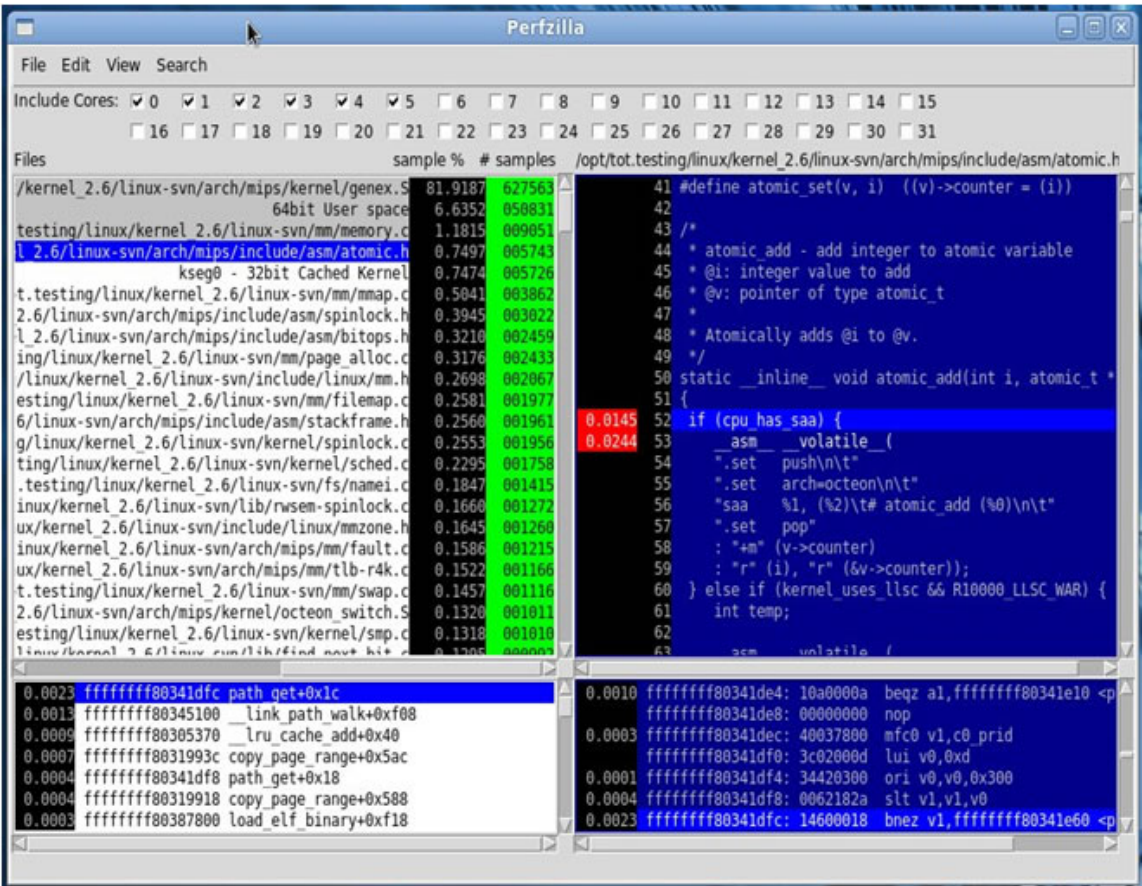
- Create a temporary file `perfzilla_objdump_elf_file.tmp` (the `objdump -d` of the `elf_file`).
- Read the profile information from the supplied `profile.log` file.

After these steps are complete, the perfzilla GUI is loaded.

## 7.4 The User Interface

The perfzilla Graphical User Interface is illustrated below:

Figure 7–1 perfzilla GUI



The perfzilla GUI includes the following elements:

- [Main Menu](#)
- [Include Cores Pane](#)
- [perfzilla Windows](#)

### 7.4.1 Main Menu

The **Main** menu includes the following items:

Menu	Option
File	New Window—Opens a new instance of perfzilla using the same ELF and profile.log files as the current instance.
	Save Source—Saves the source code file with profiling information.



Menu	Option
Edit	<b>Copy</b> —Copies the text from any of the four windows. This text can be pasted into a text editor or into an email.
View	<b>Previous Hit</b> —Moves the cursor to the previous sample information in source file.
	<b>Next Hit</b> —Moves the cursor to the next sample information in the source code.
	<b>Function Profile Info</b> —Displays the detailed profiling information.
Search	<b>Find</b> —Searches (greps) for a matching string in the source code.
	<b>Find next</b> —Searches (greps) for the next matching string in the source code.

### 7.4.2 Include Cores Pane

The **Include Cores** pane displays the maximum number of cores available on OCTEON. In the following figure, the maximum number of cores is 32:

**Figure 7–2 Include Cores Pane**



The selected check boxes represent the number of cores that are currently being profiled. The result of total samples or sample% varies depending on the number of cores being selected or deselected in this section. This is useful to control the sampling information for a single core or subset of cores.

### 7.4.3 perfzilla Windows

The perfzilla GUI includes the following windows:

- **File List** window (upper left)
- **Source Code** window (upper right)
- **Function List** window (lower left)
- **ASM** window (lower right)

Figure 7–3 File List Window

Files	sample %	# samples
testing/linux/kernel_2.6/linux-svn/arch/mips/kernel/genex.S	81.8612	10114710
64bit User space	6.6553	822319
/opt/tot.testing/linux/kernel_2.6/linux-svn/mm/memory.c	1.1886	146863
g/linux/kernel_2.6/linux-svn/arch/mips/include/asm/atomic.h	0.7475	092363
kseg0 - 32bit Cached Kernel	0.7336	090641
/opt/tot.testing/linux/kernel_2.6/linux-svn/mm/mmap.c	0.5249	064852
linux/kernel_2.6/linux-svn/arch/mips/include/asm/spinlock.h	0.3990	049303
g/linux/kernel_2.6/linux-svn/arch/mips/include/asm/bitops.h	0.3322	041052
/opt/tot.testing/linux/kernel_2.6/linux-svn/mm/page_alloc.c	0.3175	039229
nux/kernel_2.6/linux-svn/arch/mips/include/asm/stackframe.h	0.2649	032735
t/tot.testing/linux/kernel_2.6/linux-svn/include/linux/mm.h	0.2643	032656
pt/tot.testing/linux/kernel_2.6/linux-svn/kernel/spinlock.c	0.2562	031657
/opt/tot.testing/linux/kernel_2.6/linux-svn/mm/filemap.c	0.2547	031472
/opt/tot.testing/linux/kernel_2.6/linux-svn/kernel/sched.c	0.2286	028244
/opt/tot.testing/linux/kernel_2.6/linux-svn/fs/namei.c	0.1917	023690
tot.testing/linux/kernel_2.6/linux-svn/arch/mips/mm/fault.c	0.1635	020201
t.testing/linux/kernel_2.6/linux-svn/include/linux/mmzone.h	0.1622	020041
tot.testing/linux/kernel_2.6/linux-svn/lib/rwsem-spinlock.c	0.1595	019707
t.testing/linux/kernel_2.6/linux-svn/arch/mips/mm/tlb-r4k.c	0.1545	019089
/opt/tot.testing/linux/kernel_2.6/linux-svn/mm/swap.c	0.1429	017657
linux/kernel_2.6/linux-svn/arch/mips/kernel/octeon_switch.S	0.1357	016761
/opt/tot.testing/linux/kernel_2.6/linux-svn/mm/slab.c	0.1322	016337
/tot.testing/linux/kernel_2.6/linux-svn/lib/find_next_bit.c	0.1284	015863
/opt/tot.testing/linux/kernel_2.6/linux-svn/kernel/smp.c	0.1268	015670
/linux/kernel_2.6/linux-svn/arch/mips/include/asm/pgtable.h	0.1188	014684
/tot.testing/linux/kernel_2.6/linux-svn/kernel/sched_fair.c	0.1093	013504
x/kernel_2.6/linux-svn/arch/mips/cavium-octeon/octeon-irq.c	0.1042	012870
/opt/tot.testing/linux/kernel_2.6/linux-svn/lib/prio_tree.c	0.1034	012773
TLB Handler 1	0.1031	012735
opt/tot.testing/linux/kernel_2.6/linux-svn/lib/radix-tree.c	0.1021	012619

The **File List** window displays:

- A list of source files whose functions had at least one sample registered during profiling.
- The sample% and sum total of all samples collected against a particular file.
- The absolute file name path.
- Horizontal and vertical scroll bars to make it easy to navigate through the contents.

When you select a file in the **File List** window, the source code for the file displays in the **Source Code** window. The **Sample%** (in red) displays opposite a particular instruction or C source code line.

Figure 7–4 Source Code Window

```
opt/tot.testing/linux/kernel_2.6/linux-svn/fs/namei.c
114 * checking and hopefully speeding things up, we copy filenames t
115 * kernel data space before using them..
116 *
117 * POSIX.1 2.4: an empty pathname is invalid (ENOENT).
118 * PATH_MAX includes the nul terminator --RR.
119 */
120 static int do_getname(const char __user *filename, char *page)
121 {
122     int retval;
123     unsigned long len = PATH_MAX;
124
0.0002 125     if (!segment_eq(get_fs(), KERNEL_DS)) {
0.0003 126         if ((unsigned long) filename >= TASK_SIZE)
127             return -EFAULT;
0.0013 128         if (TASK_SIZE - (unsigned long) filename < PATH_MAX)
129             len = TASK_SIZE - (unsigned long) filename;
130     }
131
0.0001 132     retval = strncpy_from_user(page, filename, len);
0.0001 133     if (retval > 0) {
0.0004 134         if (retval < len)
135             return 0;
136         return -ENAMETOOLONG;
137     } else if (!retval)
138         retval = -ENOENT;
139     return retval;
140 }
141
142 char * getname(const char __user * filename)
0.0014 143 {
```

You can take the following actions from the **Source Code** window:

View samples in the file:	<ul style="list-style-type: none"><li>• Double-click the file name in the <b>File List</b> window</li><li>• Select <b>View &gt; Next</b> hit (ALT+Down in <b>Source Code</b> window)</li><li>• Select <b>View &gt; Previous</b> hit (ALT+Up in <b>Source Code</b> window)</li></ul>
Search for a pattern in the source file:	Select <b>Search &gt; Find</b> or <b>Search &gt; Find Next</b> (F5 in <b>Source Code</b> window). The search pattern is highlighted in yellow.
Copy and paste source code into a text editor or email:	Select <b>Edit &gt; Copy</b> (CTRL+C), right-click, and select <b>Paste</b> from the context menu.

Figure 7–5 Function List Window

```
0.0010 ffffffff80342bfc getname+0x5c
0.0001 ffffffff80342c08 getname+0x68
0.0001 ffffffff80342bf8 getname+0x58
```

The **Function List** window displays:

- The PC value and function name of the C source code line selected in the **Source Code** window.

**Note:** This value may not be 100% accurate.

- The sample% against the PC values and routine name.

**Figure 7–6 ASM Window**



```

0.0001 ffffffff80342be4: 14600004 bnez v1,ffffff80342bf8 <getname+0x58>
ffffff80342be8: 0050402f dsubu a4,v0,s0
ffffff80342bec: 080d0b0f j fffffff80342c3c <getname+0x9c>
ffffff80342bf0: 2410fff2 li s0,-14
ffffff80342bf4: 00000000 nop
0.0001 fffffff80342bf8: 2d021000 sltiu v0,a4,4096
0.0010 fffffff80342bfc: 14400003 bnez v0,ffffff80342c0c <getname+0x6c>

```

The **ASM** window includes the following information:

- Displays the disassembly of the instruction or C code line selected in the **Source Code** window.
- Updates the window with the sample% count and disassembly of the C instruction.

The display may span multiple lines depending on the C code. The instructions may not be continuous in some cases.

**Note:** The sample% and line number shown in the C Source code are approximate, and should not be considered 100% accurate, as the Compiler optimizations will rearrange instructions to avoid bubbles or stalls. Also note that if an instruction stalls, the sample% or count is calculated up to the line following the stalled instruction, instead of the instruction that actually stalled.

# Chapter 8

## SE-Dump and the oct-remote-coredump Tool

### 8.1 Introduction

SE-Dump and the oct-remote-coredump tool together enable the creation of a core dump file when an SE application crashes. This core file can then be analyzed using standard GDB commands.

SE-Dump is a new mechanism through which the exception handler creates an ELF core header in memory (reserved while initializing). Host/Linux side tools are used to fetch the memory location and dump the memory contents into an ELF core file.

An ELF core file is an ELF object file with a current snapshot of the memory contained in the file. Depending on the memory consumed by the application, this file can grow enormously. With an ELF core file and GDB, the user can debug the problem with source code. The available debug view is greatly improved with stack trace availability, and to view the values of global/local variables at each level of function call sequence.

The oct-remote-coredump tool is a helper tool written to coordinate with the core dumping facility built into Simple Executive. Once the SE application runs into an exception, the default behavior is to create a core dump of the program. The user can then dump the entire snapshot of the program's memory to analyze offline to discover where the error occurred.

### 8.2 Features

The oct-remote-coredump tool includes the following features:

- Capable of dumping multi-core application.
- Creates a fully GDB- capable core dump binary for analysis.
- Capable of using all protocols supported by `remote-lib` (PCI, JTAG, and so on).

### 8.3 How the Tool Works

When an exception occurs, the exception handler installed during initialization is called. In the exception handler, `cvmx_coredump_start_dump()` is called. This calculates the address at which the ELF header has to be created. The ELF header address depends on the core number. The macros defined in `cvmx-coredump.h` and `cvmx-coredump-elf.h` specify that the memory location is used to dump the ELF header along with related information.

1. The SE application's exception handler calculates the offset for saving registers.
2. SE-Dump starts to build an ELF program header, for which it needs the TLB entries. Through the TLB entries, it determines code/data and stack information for populating the ELF header.
3. Once the TLB and registers are in place, dumping is started.
4. The ELF header is created.

5. A program header is created with one PT\_NOTE section and the value of (the number of TLBs + the number of memory allocations).

Once the program headers are in place, `oct-remote-coredump` runs to complete the process. The `oct-remote-coredump` tool reads the ELF header along with program headers. First, the ELF header is written, then the program headers, and then the PT\_NOTE section. Finally, for each entry in the program header, it copies the target's memory to file.

## 8.4 Command-line Option

The supported command-line option for `oct-remote-coredump` is:

```
oct-remote-coredump <coremask>
```

where `<coremask>` is the mask used to run the program. If no mask was provided (using `boot-oct`), `0x1` should be used. This argument is mandatory.

## 8.5 Getting Started

The sources for the `oct-remote-coredump` utility are located in the `${OCTEON_ROOT}/host/coredump-util` directory. SE-Dump is enabled by default in Simple Executive applications, so no action is required to enable it.

The following steps describe how to use SE-Dump and `oct-remote-coredump` to generate core dump files for analysis with GDB:

### 1. Build the SE application.

Build your own SE application, or use the sample program provided at `<OCTEON_ROOT>/examples/coredump` for trying out the tool.

To build the sample program (`coredump`), run the following commands:

```
$ cd <OCTEON_ROOT>/examples/coredump
$ make
```

### 2. Load and run the application.

Transfer the resulting binary from step 2 to the target (for details on loading the application, see the applicable sections of the *OCTEON® Programmer's Guide, The Fundamentals* (Curtis, 2009). At the U-boot prompt, run the command:

```
Octeon-ebb5200# bootoct $(loadaddr) coremask=0x3 endbootargs
```

The sample program is designed to crash after calculating some checksums of SHA1, MD5, and Kasumi. Once the program crashes, you can run the `oct-remote-coredump` tool.

### 3. Dump the corefile.

Run the following command (from either the host or on a target running Linux):

```
$ oct-remote-coredump 0x3.
```

This creates sequentially named files —`CVMX_DUMP.CORE_0` and `CVMX_DUMP.CORE_1`—for `core0` and `core1` (which is specified in the command line by `coremask 0x3`).

### 4. Run gdb to analyze the crash global/local variables.

Because each core has different data/stack segments allocated while running the application, it is impossible to encode all the information related to the application in



the ELF file format. Consequently, separate core files are generated for each core. You must repeat this step for each core.

An example GDB session is shown below.

**Table 8–1 gdb session**

```
$ mips64-octeon-linux-gnu-gdb
<<snip>>
(gdb) symbol-file coredump-crypto
Reading symbols from examples/coredump/coredump-crypto...done.
(gdb) core-file CVMX_DUMP.CORE_0
[New process 1]
#0 0x0000000010000d94 in crash () at coredump-crypto.c:745
745          *p = 0x1234;
(gdb) backtrace
#0 0x0000000010000d94 in crash () at coredump-crypto.c:745
#1 0x0000000010001294 in test_snow3g () at coredump-crypto.c:796
#2 test_kasumi () at coredump-crypto.c:829
#3 test_aes () at coredump-crypto.c:861
#4 0x0000000000000000 in ?? ()
(gdb) info registers
Id Target Id Frame
* 1 process 1 0x0000000010000d94 in crash () at coredump-crypto.c:745

R0      0000000000000000 zero      at      v0      0000000000000000 v1
          a0      ffffffff bffffff 0000000000001234 000000000000000a
R4      ffffffff ffffffff a0      a1      a2      a3
          a4      000000000000000a 0000000000000019 0000000000000019
          a5      a6      a7
R8      ffffffff 7f0000 0a0a0a0a0a0a0a0a 000000000000002a 0000000000000000
          t0      t1      t2      t3
R12     0000000000000012 0000000000000000 0000000000000000 0000000000000000
          s0      s1      s2      s3
R16     0000000012016150 0000000000000001 00000000101c9328 ffffffff c0784fb0
          s4      s5      s6      s7
R20     0000000000000028 0000000000000000 0000000000000000 0000000000000000
          t8      t9      k0      k1
R24     0000000000000000 0000000010000238 0000000010034748 ffffffff 80001090
          gp      sp      s8      ra
R28     00000000120091d0 00000000160fff98 0000000016100000 0000000010000784
          sr      lo      hi      bad
          00000000501004e3 0000000000000000 0000000000000000 0000000000000000
          cause pc
          000000004000800c 0000000010000d94
          fsr      fir
gdb-script0:5: Error in sourced command file:
can't read register 160 (fsr)
(gdb) print plaintext
$1 = 0xi01c9328 "Original plaintext message"
(gdb)
```

# Chapter 9

## The viewzilla Tool

### 9.1 Introduction

The OCTEON SDK includes the performance profiling tool viewzilla, which runs on the OCTEON hardware simulator. The tool collects profiling data and provides a visual interface for interpreting that data, showing the application programmer where the code can be optimized for OCTEON.

**Important!** viewzilla only runs on the hardware simulator; it does not run on actual hardware.

### 9.2 Features

The viewzilla tool includes the following features:

- Provides a visual representation of each processor's pipeline.
- Handles detailed load for each processor's crypto engine (`-modes=cryptovz` for the simulator).
- Manages Level 1 and Level 2 cache operations.
- Controls coherent bus utilization.
- Utilizes low latency memory bandwidth.
- Displays C Source.
- Performs instruction tracing per processor or for all processors.
- Enables fast searching.
- Provides loop visualization with instruction stepping.
- Facilitates quick scale view from high-level overview down to individual instructions.

### 9.3 Getting Started

The viewzilla tool requires the following two files as input:

- `octeon.vz`—the viewzilla file to load. The file is generated by running `oct-sim` with the `-vz` option. The option `-modes=cryptovz` can also be used to detail the crypto units.
- `elf`—Optional Mips ELF files from which to load debugging information.

To start viewzilla, run the following command:

```
$ viewzilla octeon.vz [elf] [elf] ...
```

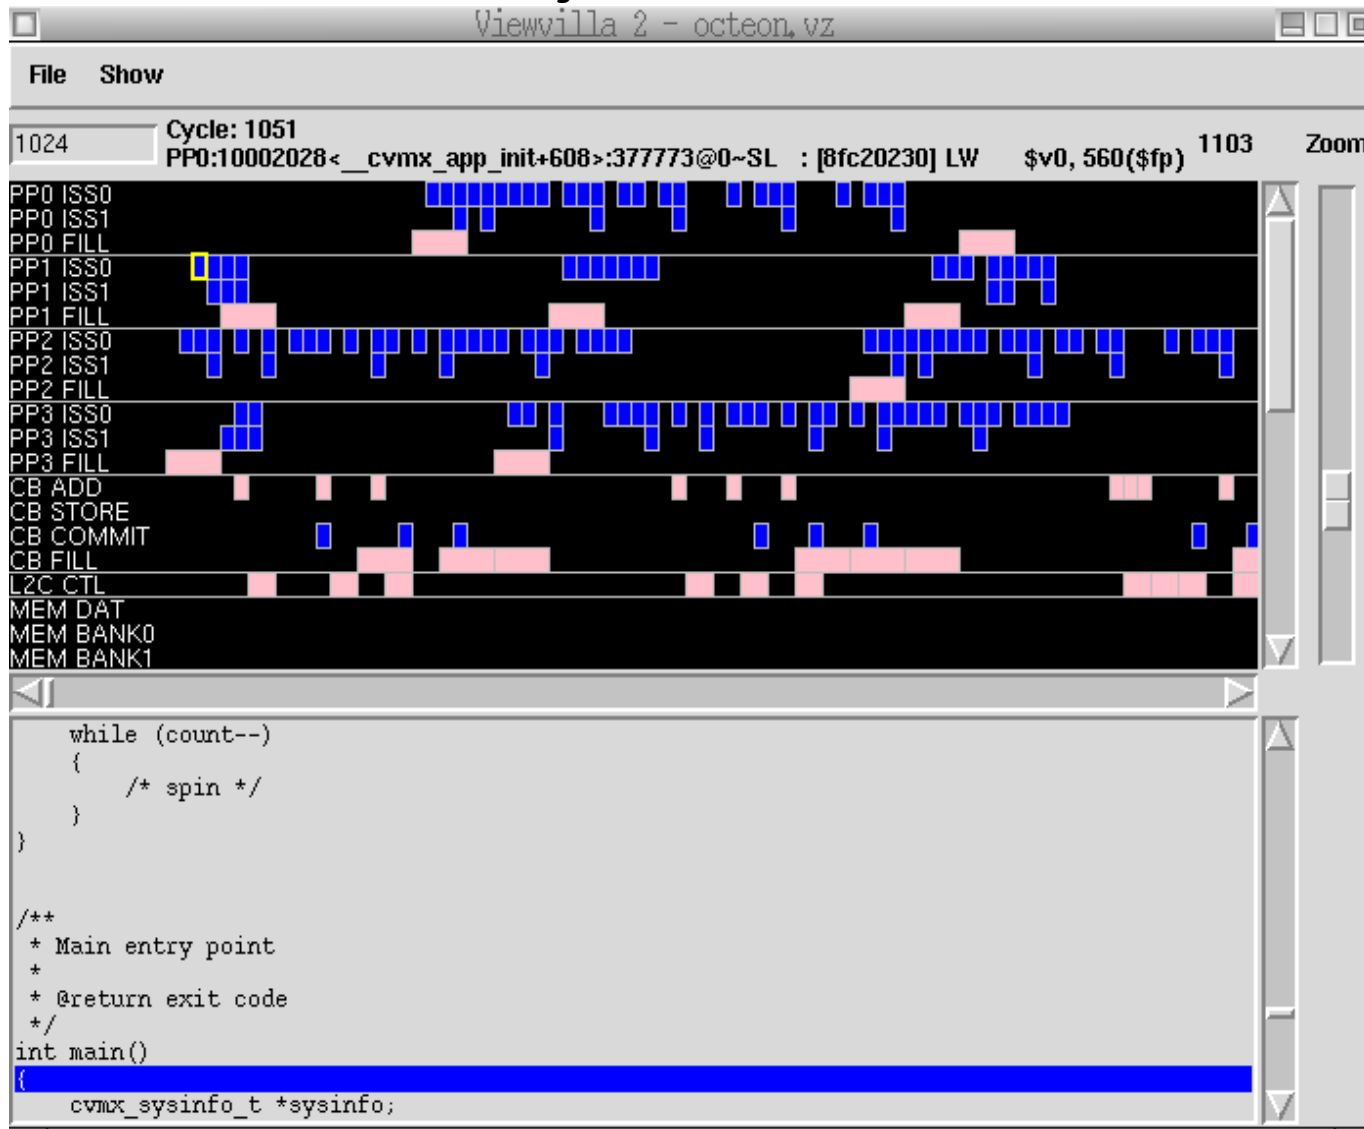
To use viewzilla, build applications with DWARF debugging information (`-g` compiler option).



## 9.4 The User Interface

The viewzilla Graphical User Interface is illustrated below:

**Figure 9–1 viewzilla GUI**



The viewzilla GUI includes the following elements:

- [Main Menu](#)
- [Windows and Navigation Tools](#)

### 9.4.1 Main Menu

The **Main** menu includes the following items:

**Table 9–1 viewzilla Main Menu**

Menu	Option
File	<b>Find</b> — Used to find arbitrary strings in the viewzilla file. Searching starts at the current cycle and continues forward through time. Searches are case sensitive. Short cut keys: <b>Control-F</b>
	<b>Find Next</b> — Continues the previous search. Short cut keys: <b>Control-G, n.</b>
	<b>Quit</b> — Exits the viewzilla program.
Show	<b>Source</b> —Uses the DWARF debugging information from the ELF files to show the original program source.
	<b>Processor Instructions</b> —Displays the instruction trace annotated by source for the currently selected processor.
	<b>All Processor Instructions</b> —Displays the instruction traces annotated by source for all the running processors.
	<b>Highlight Next</b> —Moves the highlight to the next active cycle for the selected hardware unit. Short cut keys: <b>-&gt;, k.</b>
	<b>Highlight Previous</b> —Moves the highlight to the previous active cycle for the selected hardware unit. Short cut keys: <b>&lt;-, j.</b>
	<b>Hardware</b> —Controls whether each hardware unit is displayed.

### 9.4.2 Windows and Navigation Tools

Other windows and navigation tools on the viewzilla interface are described in the following table:

**Table 9–2 viewzilla Windows and Navigation Tools**

Menu	Option
Start cycle input field	The first clock cycle displayed at the left edge of the cycle view window. This value may be modified by directly typing a value, scrolling the horizontal scrollbar, or by using one of the searching operations.
Selected box description	Displays a short description of what is occurring whenever the mouse moves over an active cycle.
Ending cycle	The last cycle visible in the cycle view window. It is calculated based on the current zoom level and window width.
Cycle view window	Displays each hardware unit (vertical axis) and its state at each cycle (horizontal axis). Moving the mouse over boxes shows a short description. Click a box to select it and any boxes associated with it; available source/instructions are displayed. Step to the next/previous selection using the left and right arrow keys.
Vertical scrollbar for the cycle view window	If the viewzilla window is too small, not all hardware units can be display in the cycle view window. Use this scrollbar to control which ones are available.
Zoom level slider	Controls how much information is displayed. Slide up to show fewer cycles, but more detail. Slide down to show more cycles, but less detail.

Menu	Option
Horizontal scrollbar for the cycle view window	Quickly move to any cycle in the viewzilla trace. The left and right arrows move one cycle at a time. Click the page back/forward arrows to move 3/4s of the view area at a time.
Source / Instruction display window	Depending on the selection in the <b>Show</b> pulldown menu, displays either the source file for the current instruction or an instruction trace with source annotation. <b>Note:</b> This is only available when an instruction is selected.
Vertical scrollbar for the source window	Scrolls through the loaded source file or instruction trace.

## 9.5 Understanding Output

When you click a processor instruction in the viewzilla interface to select it, text detailing the state of the processor at execution time displays. The following examples explain how to decode the text.

### Example 1:

```
PP0:20001254<main+116>:2400@0~MisP : [24a50001] ADDIU $a1, $a1, 1
[$a1=2]
```

- PP0—The core to which the information applies; in this case, core 0
- 20001254<main+116>—The instruction virtual address (hex + label)
- 2400@0—Cycle 2400, issued on pipe 0
- ~MisP—viewzilla code indicating that the instruction was delayed due to a branch mispredict
- 24a50001—Hex of the instruction
- ADDIU \$a1, \$a1, 1—Instruction mnemonic

### Example 2:

```
PP0:20001264<main+132>:7302@0~DWB : [dfbf0000] LD $ra, 0($sp)
[$ra=20000264(7FFFFFFD0,7FFFFFFD0)]
```

- ~DWB—viewzilla code indicating that this load missed and matched a pending write buffer
- (7FFFFFFD0,7FFFFFFD0)—(virtual address,physical address) in hex
- \$ra=20000264—0x20000264 was loaded from physical address 0x7FFFFFFD0

### Example 3:

```
PP0:200004EC<flush_print+ 36>:7383@0~SL : [ffb00000] SD $s0,
0($sp) [(7FFFFFFD0,7FFFFFFD0)=0]
```

- ~SL—viewzilla code indicating that the instruction was delayed because it had to issue on pipe 0
- (7FFFFFFD0,7FFFFFFD0)—(virtual address,physical address) in hex
- =0—The value 0 was written to physical address 0x7FFFFFFD0

### 9.5.1 Delay/Stall Codes

The following table provides a list of codes that viewzilla can return due to delays or stalls:

**Table 9–3 viewzilla: List of Delay/Stall Codes**

Stall	Reasons
~ITLB	Instruction TLB miss
~BR	Prior branch instruction
~JR	Prior jump instruction
~MisP	Prior branch mispredicted
~MTCOP0	Prior MT COP0 instruction
~IF	Instruction fetch
~SL	Needs pipe 0, but pipe 1 is available
~BRF	Ignore this
~RAW	Input register values not available
~MisJR	Mispredicted jump
~WAI	WAIT instruction
~MUL	Multiplier busy
~CR	Crypto busy
~LLM	LLM busy
~D	Dcache miss that does not hit a write buffer entry
~DWB	Dcache miss that hits a write buffer entry
~SY	SYNC/SYNCS/SYNCIOBDMA instruction

## 9.6 The vz-cut Utility

The `vz-cut` utility provides an easy way to copy a section of a viewzilla file into a smaller file. The original file is not modified.

To start `vz-cut`, run the following command:

```
$ vz-cut --start start_cycle --end end_cycle [--filename
input_filename] [--outfilename output_filename] [--help]
```

where

- `--start`—The cycle on which to start copying; defaults to 0
- `--end`—The cycle on which to end copying; defaults to the end of the file
- `--filename`—Input filename; defaults to `octeon.vz`
- `--outfilename`—Output filename; defaults to `'octeon-cut.vz'`
- `--help`—Displays help message.

**Note:** Only one of `start` or `end` is required because `start` defaults to 0 and `end` defaults to the end of the file if the other value is not provided.

## 9.7 Tips for Running viewzilla

The following list includes tips for running viewzilla effectively:

- Instruction tracing is already contained in the viewzilla file. To suppress tracing, run `oct-sim` with the `-quiet` option.
- Most of the time you are only interested in application performance. Use the option `-wait=main` to skip all analysis before the main entry point of the program. Without this option, it can be difficult to wade through all the output for the bootloader and C library startup.
- The simulator runs much more slowly with viewzilla output turned on. Use the `-wait` option to start output only in the area in which you are interested.
- viewzilla output files can get large very quickly. Make sure you have adequate free disk space. As an example, booting 16 cores and running for 10 million cycles generates a 14 GB output file.
- Compile applications to be traced with debugging symbols (`-g`).
- viewzilla uses the full path names present in the DWARF debugging information. Make sure source code can be found at these locations. If the source view is unable to find a file, it displays the path where it tried to locate the source.
- Due to the large amount of data, searching through a viewzilla trace can be slow. Avoid long searches by scrolling to the approximate cycle where your search result is expected to be found.

# Appendix A

## Controlling SE Applications from Linux

This Appendix provides instructions for using the `oct-app-ctl` utility, included in the SDK for controlling Simple Executive (SE) applications from Linux. Potential uses for controlling standalone SE applications from Linux include:

- Packaging the entire application, including control and data plane, in a single embedded root file system, or enabling a "live" upgrade of the data plane portion without rebooting the control plane.
- Dynamic load balancing where the number of CPU cores can be switched between the data plane SE application and Linux to match changing load patterns dynamically. In such cases, the dynamic load-control supervisor (not provided) would employ the Linux CPU hot-plugging mechanism to make CPU cores available from Linux and added to SE applications and vice versa.

**Note:** This utility replaces the `bootoct` and `oct-shutdown-app` utilities included in previous releases of the SDK.

### A.1 Using the `oct-app-ctl` Utility

The `oct-app-ctl` utility supports five operations, described in the following sections:

- [Starting an SE Application](#)
- [Displaying Application Status](#)
- [Shutting Down an Application](#)
- [Adding Cores to an Application](#)
- [Deleting Cores from an Application](#)

For each operation, the exact action can be specified either as an argument to `oct-app-ctl` or by using a link to the utility executable identifying the operation as follows:

```
oct-app-ctl <operation>
```

or

```
oct-app-<operation>
```

#### A.1.1 Starting an SE Application

To start a SE application, run one of the following commands:

```
oct-app-ctl boot <options> <app_file_name> [<app_arguments>]  
oct-app-boot <options> <app_file_name> [<app_arguments>]
```

The supported command-line options for `oct-app-ctl` when booting an application are listed in the following table.

**Table A-1 Command-line Options for Booting an Application**

Option	Description
-verbose=<num>	Increase verbosity level of the program up to the value of 3. Default is 0.
-numcores=<N>	The number of CPU cores on which to start the application. Default is 1.
-mask=<hex-number>	The core mask of CPUs on which to start the application.
-heapsize=<bytes>	Set the size of the heap memory allocated per CPU core. Default is 3 MB.
-stacksize=<bytes>	Set the size of the stack memory allocated per CPU core. Default is 1 MB.

Note the following when starting applications:

- This method can be used to start any SE application, whether or not the application supports hot-plugging (the remaining operations require the support of hot-plugging).
- The core mask argument (mask=<hex-number>) only supports up to 64 cores. Consequently, if the core mask argument is specified and includes cores that are not available, the unavailable cores are replaced with the same number of available cores to satisfy the core count in the requested mask, if possible.

### A.1.2 Displaying Application Status

To display the status of applications, run one of the following commands:

```
oct-app-ctl info
oct-app-info
```

The command returns the following information:

- A list of which SE applications are currently running on which cores.  
**Note:** The cores currently being used by Linux are not specified. In order to control and display the CPU cores running Linux, refer to the Linux CPU hot-plugging documentation.
- Active core mask—Applications that are capable of supporting hot-plugging, which is required for shutting down applications and adding/removing cores.
- An index entry for every instance of an application. This information is useful if there are multiple applications with the same name (because they were started from the same executable file) and you wish to perform an operation on a specific instance of the application.
- The core mask of currently unused CPU cores that are available for starting new applications or adding into existing applications.

### A.1.3 Shutting Down an Application

To shut down an application on all CPU cores on which it is running, run one of the following commands:

```
oct-app-ctl shutdown <options> {-index=<app-index> | <app-name>}
oct-app-shutdown <options> {-index=<app-index> | <app-name>}
```

The following option is available (it is not required):

**Table A-2 Command-line Options for Shutting Down Applications**

Option	Description
<code>-verbose=&lt;num&gt;</code>	Increase verbosity level of the program up to the value of 3. Default is 0.

Note the following when shutting down an application:

- The application being stopped must support hot-plugging and must have call-back routines installed so that it can be notified of requests for relinquishing or adding resources. See the example program located at `sdk/examples/hotplug_app` for additional information.
- The application affected must be identified either by its name, which is identical to the name of the executable object file that was used to start the application, or by means of the application instance index number.
- The executable file does not need to be present in the current (or any other) directory for this operation to succeed, as it is not accessed for anything other than the `boot` operation.

### A.1.4 Adding Cores to an Application

To add cores to an application, run one of the following commands:

```
oct-app-ctl add <options> {-index= <app-index> | <app-name>}
oct-app-add <options> {-index= <app-index> | <app-name>}
```

The following table lists the options available when adding cores.

**Table A-3 Command-line Options for Adding Cores**

Option	Description
<code>-verbose=&lt;num&gt;</code>	Increase verbosity level of the program up to the value of 3. Default is 0.
<code>-numcores=&lt;N&gt;</code>	The number of CPU cores on which to start the application. Default is 1.
<code>-mask=&lt;hex-number&gt;</code>	The core mask of CPUs on which the application should be started.

Note the following when adding cores:

- The application must support hot-plugging and must have call-back routines installed so that it can be notified of requests for relinquishing or adding resources. See the example program located at `sdk/examples/hotplug_app` for additional information.
- The core mask argument (`mask=<hex-number>`) only supports up to 64 cores; it is recommended that the `-numcores` argument be used instead.
- If the core mask argument is specified and includes cores that are not presently available, the unavailable cores are replaced with the same number of available cores to satisfy the core count in the requested mask, if possible.

### A.1.5 Deleting Cores from an Application

To delete cores from an application, run one of the following commands:

```
oct-app-ctl del <options> {-index= <app-index> | <app-name>}
oct-app-del <options> {-index= <app-index> | <app-name>}
```



The following table lists the options available when deleting cores.

**Table A-4 Command-line Options for Deleting Cores**

Option	Description
<code>-verbose=&lt;num&gt;</code>	Increase verbosity level of the program up to the value of 3. Default is 0.
<code>-numcores=&lt;N&gt;</code>	The number of CPU cores to relinquish. Default is 1.
<code>-mask=&lt;hex-number&gt;</code>	The coremask of CPUs that should be stopped.

Note the following when deleting cores:

- The application must support hot-plugging and must have call-back routines installed so that it can be notified of requests for relinquishing or adding resources. See the example program located at `sdk/examples/hotplug_app` for additional information.
- This operation will refuse to delete the initial core running the application, which is the lowest-numbered core on which the application was initially started. All other cores can be deleted, including cores that are numerically lower than the initial core but were added using the add operation after the application was started.
- If a core mask is specified in the command line that includes cores not running an instance of the application, or if the mask specifies the initial core, the actual mask is modified to stop the same number of cores as requested in the mask, if possible.

## A.2 Limitations

The `oct-app-ctl` utility has the following limitations:

- The utility presently does not support bare-metal applications that have been started from the bootloader. It cannot add cores, remove cores, or shut such applications down. If deploying Linux, it is recommended that the bootloader only be used to start the kernel, with any further control over bare-metal SE applications performed from Linux.
- The utility has only been tested with big-endian executables and hosts. It does not currently support little-endian versions of either.

# Appendix B

## Examples

### B.1 oct-profile

#### B.1.1 Invoking oct-profile for traffic-gen

oct-profile invoked for traffic-gen (over PCIe).

**Table B-1 oct-profile: Invoking**

Oct-profile /tftpboot/traffic-gen						
core#:	0	1	2	3	4	5
2356	0	0	0	0	0	main
2356	0	0	0	0	0	statistics_gatherer
0	3184	3184	3184	3184	3184	packet_transmitter
0	162	166	138	157	161	cvmx_get_cycle
0	162	166	138	157	161	cvmx_clock_get_count
0	14	17	21	22	20	cvmx_send_single
0	14	17	21	22	20	cvmx_write64_uint64
0	0	0	0	0	170	cvmx_scratch_read64
0	137	116	127	121	146	cvmx_scratch_write64
489	0	0	0	0	0	periodic_update
20	0	0	0	0	0	cvmx_get_cycle
20	0	0	0	0	0	cvmx_clock_get_count
339	0	0	0	0	0	uart_printf

#### B.1.2 Profiling traffic-gen over PCI

Profile traffic-gen running over PCI; Display PC along with sample hits associated with routines.

**Table B-2 oct-profile: Profiling traffic-gen over PCI**

oct-profile -address /tftpboot/traffic-gen						
core#:	0	1	2	3	4	5
1	0	0	0	0	0	__sfvwrite_r
1	0	0	0	0	0	0x00000000101418a0
2	0	0	0	0	0	vfprintf_r
1	0	0	0	0	0	0x000000001013bfe0
1	0	0	0	0	0	0x000000001013c610
2719	0	0	0	0	0	main
2719	0	0	0	0	0	statistics_gatherer
17	0	0	0	0	0	0x0000000010028730
13	0	0	0	0	0	0x000000001002873c
18	0	0	0	0	0	0x0000000010028740
24	0	0	0	0	0	0x0000000010028744
81	0	0	0	0	0	0x0000000010028748
23	0	0	0	0	0	0x0000000010028750
860	0	0	0	0	0	0x0000000010028754
5	0	0	0	0	0	0x000000001002875c
119	0	0	0	0	0	0x0000000010028764
146	0	0	0	0	0	0x000000001002876c
24	0	0	0	0	0	0x000000001002878c
32	0	0	0	0	0	0x0000000010028790
2	0	0	0	0	0	0x0000000010028798
4	0	0	0	0	0	0x000000001002879c
7	0	0	0	0	0	0x00000000100287a0
1344	0	0	0	0	0	0x00000000100287a4
0	3636	3636	3636	3636	3636	packet_transmitter
0	100	117	120	129	220	0x0000000010009410
0	90	88	99	96	176	0x0000000010009418
0	268	259	269	289	494	0x000000001000941c
0	95	105	109	82	165	0x000000001000961c
0	378	297	278	256	492	0x0000000010009624
0	83	117	106	90	192	0x0000000010009628
0	293	287	299	258	554	0x000000001000962c
0	87	95	105	94	162	0x0000000010009634

0	287	294	276	305	515	0x000000001000971c
0	15	22	19	23	26	0x0000000010009728
0	263	263	300	266	0	0x0000000010009760
0	297	301	267	281	0	0x0000000010009764
0	87	75	87	89	0	0x000000001000976c
0	264	254	262	294	0	0x0000000010009770
0	95	81	77	88	0	0x0000000010009e94
0	105	107	97	101	0	0x0000000010009e9c
0	185	189	182	184	0	0x0000000010009ea0
0	244	251	224	243	0	0x0000000010009ea4
0	191	181	187	193	170	cvmx_get_cycle
0	191	181	187	193	170	cvmx_clock_get_count
0	96	94	94	105	170	0x0000000010009618
0	95	87	93	88	0	0x0000000010009e90
0	81	97	113	136	108	cvmx_pow_work_request_async_nocheck
0	46	53	57	85	70	0x0000000010009738
0	20	24	32	26	15	0x000000001000973c
0	15	20	24	25	23	cvmx_send_single
0	15	20	24	25	23	cvmx_write64_uint64
0	15	20	24	25	23	0x0000000010009740
0	0	0	0	0	184	cvmx_scratch_read64
0	0	0	0	0	184	0x0000000010009718
0	128	156	160	139	178	cvmx_scratch_write64
0	128	156	160	139	178	0x0000000010009724
573	0	0	0	0	0	periodic_update
16	0	0	0	0	0	0x000000001001aaf0
8	0	0	0	0	0	0x000000001001aaf4
9	0	0	0	0	0	0x000000001001aaf8
14	0	0	0	0	0	0x000000001001aafc
19	0	0	0	0	0	0x000000001001ab00
13	0	0	0	0	0	0x000000001001ab04
9	0	0	0	0	0	0x000000001001ab08
21	0	0	0	0	0	0x000000001001ad28
35	0	0	0	0	0	0x000000001001ad2c
28	0	0	0	0	0	0x000000001001ad30
1	0	0	0	0	0	0x0000000010023a04
34	0	0	0	0	0	cvmx_get_cycle
34	0	0	0	0	0	cvmx_clock_get_count
34	0	0	0	0	0	0x000000001001ab20
341	0	0	0	0	0	uart_printf
5	0	0	0	0	0	0x0000000010004f60
1	0	0	0	0	0	0x0000000010004f9c
1	0	0	0	0	0	0x0000000010004fb8
170	0	0	0	0	0	0x0000000010004fd0

## B.2 oct-remote-profile

### B.2.1 Profiling traffic-gen

Profiling traffic-gen application running on OCTEON connected over PCIe.

**Table B-3 oct-remote-profile: Profiling traffic-gen**

```

$oct-remote-profile

```

core#:	0	1	2	3	4	5
0x0000000010004f60:	4	0	0	0	0	0
0x0000000010004fd0:	170	0	0	0	0	0
0x0000000010004fd4:	87	0	0	0	0	0
0x0000000010004fd8:	74	0	0	0	0	0
0x0000000010005008:	4	0	0	0	0	0
0x0000000010005010:	2	0	0	0	0	0
0x0000000010009410:	0	99	84	104	128	194
0x0000000010009418:	0	101	88	99	83	168
0x000000001000941c:	0	222	231	230	234	456
0x0000000010009618:	0	65	84	86	95	59
0x000000001000961c:	0	69	80	99	99	158
0x0000000010009624:	0	334	277	245	265	469
0x0000000010009628:	0	82	111	112	90	140
0x000000001000962c:	0	242	253	262	245	507
0x0000000010009634:	0	85	82	94	90	159
0x0000000010009718:	0	0	0	0	0	176
0x000000001000971c:	0	236	256	242	256	461
0x0000000010009724:	0	150	142	136	124	114
0x0000000010009728:	0	27	14	21	11	18
0x0000000010009738:	0	47	54	54	60	65
0x000000001000973c:	0	13	18	17	12	17
0x0000000010009740:	0	18	13	20	15	17
.						
0x0000000010028764:	146	0	0	0	0	0
0x000000001002876c:	173	0	0	0	0	0
0x000000001002878c:	35	0	0	0	0	0
0x0000000010028790:	41	0	0	0	0	0
0x0000000010028798:	3	0	0	0	0	0
0x000000001002879c:	2	0	0	0	0	0
0x00000000100287a0:	2	0	0	0	0	0
0x00000000100287a4:	1168	0	0	0	0	0
0x000000001003dea8:	2	0	0	0	0	0
0x000000001013bfe0:	2	0	0	0	0	0

## B.2.2 Hardware Statistics Profiling traffic-gen

Print only perf counters, OCTEON hardware statistics when Profiling traffic-gen over PCIe.

**Table B-4 oct-remote-profile: Hardware Statistics**

```

$ oct-top -s

```

Core counters:	
brmis 63407[0], 28648414[1], 28648708[2], 28649755[3], 28649762[4], 28682185[5],	
total (excluding 0) brmis 143278824,	
sync 0[0], 0[1], 0[2], 0[3], 0[4], 0[5],	
total (excluding 0) sync 0,	
L2 statistics for TAD 0:	
bus_xmc(addr) count:	645
bus_xmd(store) count:	546
bus_rsc(commit) count:	645
bus_rsd(fill) count:	644
bus_ioc(IO req) count:	644
bus_ior(IO req) count:	644
hit count:	590258
miss count:	0
lfb-wait-vab bus utilization:	0%
lfb-wait-lfb bus utilization:	0%
DRAM Statistics:	
DRAM ops count:	142357388
dclk count:	5530240452959
utilization:	0.00%
FPA Pool statistics:	
FPA Pool 0 (Packet Buffers):	816 free
FPA Pool 1 (WQE buffers):	490 free
FPA Pool 2 (PKO cmd buffers):	0 free, PKO output queues 128
POW/SSO Qos Input Queue Statistics	
Queue 0: WQE count	0, Buffers allocated 0
Queue 1: WQE count	0, Buffers allocated 0
Queue 2: WQE count	0, Buffers allocated 0
Queue 3: WQE count	0, Buffers allocated 0
Queue 4: WQE count	0, Buffers allocated 0
Queue 5: WQE count	0, Buffers allocated 0
Queue 6: WQE count	0, Buffers allocated 0
Queue 7: WQE count	0, Buffers allocated 0
Total WQEs in all QOS levels: 0 (in-unit & in-memory)	
Total Free Buffers in all QOS levels: 243	

## B.2.3 Profiling traffic-gen Over Subset of Cores

Profile traffic-gen running over PCIe on subset of cores.

**Table B-5 oct-remote-profile: Profiling traffic-gen over Subset of Cores**

\$ Oct-remote-profile --core_mask=0x23			
core#:	0	1	5
0x0000000010004f60:	1	0	0
0x0000000010004f74:	1	0	0
0x0000000010004fd0:	129	0	0
0x0000000010004fd4:	81	0	0
0x0000000010004fd8:	74	0	0
0x0000000010004fe0:	1	0	0
0x0000000010005008:	1	0	0
0x0000000010005010:	1	0	0
0x0000000010009410:	0	95	193
0x0000000010009418:	0	79	177
0x000000001000941c:	0	234	482
0x0000000010009618:	0	67	156
0x000000001000961c:	0	74	166
0x0000000010009624:	0	350	496
0x0000000010009628:	0	87	144
0x000000001000962c:	0	252	449
0x0000000010009634:	0	69	146
0x0000000010009718:	0	0	149
0x000000001000971c:	0	246	445
0x0000000010009724:	0	116	120
0x0000000010009728:	0	16	12
0x0000000010009738:	0	46	49
0x000000001000973c:	0	19	17
0x0000000010009764:	0	250	0
0x000000001000976c:	0	71	0
0x0000000010009770:	0	237	0
0x0000000010009e90:	0	80	0
.	.	.	.
0x00000000100287a0:	3	0	0
0x00000000100287a4:	1197	0	0
0x000000001003deac:	1	0	0
0x0000000010138040:	1	0	0

## B.2.4 Available Events

List of available events.**oct-remote-profile: Available events:**

\$ oct-remote-profile -l			
clk	issue	ret	nissue
sissue	dissue	ifi	br
brmis	j	jmis	replay
iuna	trap	uunload	uustore
upload	ustore	ec	mc
cc	csrc	cfetch	cpref
ica	ii	ip	cimiss
wbuf	wdat	wbufld	wbuffl
wbuftr	badd	baddl2	bfill
ddids	idids	didna	lds
lmlds	iolds	dmlds	sts
lmsts	iosts	iobdma	dtlb
dtlbad	itlb	sync	synciob
syncw	eretmis	likmis	hazard-trap

## B.2.5 Event-based Profiling for traffic-gen

Start event-based profiler for traffic-gen (running on target over PCIe) for Cache miss event and sample rate of 100000.

**Table B–6 oct-remote-profile: Event-based Profiling for traffic-gen**

```
$ oct-remote-profile -events=cimiss:100000
```

core#:	0	1	2	3	4	5
0x0000000010004f60:	3	0	0	0	0	0
0x0000000010004f74:	1	0	0	0	0	0
0x0000000010004fd0:	124	0	0	0	0	0
0x0000000010004fd4:	48	0	0	0	0	0
0x0000000010004fe0:	7	0	0	0	0	0
0x0000000010005008:	2	0	0	0	0	0
0x000000001000500c:	1	0	0	0	0	0
0x0000000010009410:	0	35	34	25	28	160
0x0000000010009418:	0	27	30	33	30	41
0x000000001000941c:	0	58	60	61	60	126
0x0000000010009618:	0	28	29	36	27	172
0x000000001000961c:	0	59	67	62	61	153
0x0000000010009624:	0	126	110	140	118	145
0x0000000010009628:	0	31	31	30	25	153
0x000000001000962c:	0	69	58	56	77	184
0x0000000010009634:	0	0	0	0	0	183
0x0000000010009718:	0	0	0	0	0	3
0x000000001000971c:	0	94	130	96	83	205
0x0000000010009724:	0	48	78	61	67	114
0x0000000010009728:	0	28	28	26	25	27
0x0000000010009738:	0	26	21	31	25	36
0x000000001000973c:	0	27	34	28	26	31
0x0000000010009760:	0	275	283	276	315	0
0x0000000010009764:	0	183	205	179	174	0
0x0000000010009770:	0	199	179	191	218	0
0x0000000010009e90:	0	105	96	105	102	0
0x0000000010009e94:	0	111	80	103	105	0
0x0000000010009e9c:	0	93	97	96	91	0
0x0000000010009ea0:	0	86	88	81	81	0
0x0000000010009ea4:	0	113	83	104	83	0
0x0000000010028754:	1186	0	0	0	0	0
0x000000001002876c:	1	0	0	0	0	0
0x00000000100287a0:	1	0	0	0	0	0
0x00000000100287a4:	1630	0	0	0	0	0
0x000000001013bfe0:	1	0	0	0	0	0
0x000000001013c5e0:	1	0	0	0	0	0

^C

```
Event      : cimiss
Sampling Rate : 1000000
Total samples : 17705
```

Core#:	0	1	2	3	4	5
Samples:	4433	2681	2680	2680	2680	2551

## B.3 oct-top

### B.3.1 Frequent Routines

List the frequently called routines (in traffic-gen) in increasing order of number of sample hits.

**Table B–7 oct-top Frequent Routines**

```
$ oct-top /tftpboot/traffic-gen
```

Function Name:	Total hits
_vfprintf_r:	1
cvmx_sysinfo_get:	1
fdt_next_tag:	2
x_scratch_read64:	6
cvmx_get_cycle:	17
cvmx_clock_get_count:	17
cvmx_send_single:	123
cvmx_write64_uint64:	123
fastpath_receive:	284
periodic_update:	396
uart_printf:	411
cvmx_scratch_write64:	758
cvmx_get_cycle:	846
cvmx_clock_get_count::	846
main:	2893
statistics_gatherer:	2893
packet_transmitter:	18375

### B.3.2 Top Information on a Core

Generate top information of an application running on a specific core.

**Table B-8 oct-top: Top Core Information**

```
$ oct-top /tftpboot/traffic-gen --cpu=0
Function Name:                Total hits
cvmx_scratch_read64:          1
cvmx_send_single:             22
cvmx_write64_uint64:          22
fastpath_receive:             59
cvmx_get_cycle:               184
cvmx_clock_get_count:         184
cvmx_scratch_write64:        189
packet_transmitter:           3571
```

### B.3.3 PC Values

List the PC values in ascending order.

**Table B-9 oct-top: PC Values in Ascending Order**

```
$ oct-top
0x0000000010004f8c:          1
0x00000000100078e0:          1
0x0000000010009a8c:          1
0x0000000010009fb4:          1
0x00000000100287e0:          1
0x0000000010140408:          1
0x0000000010005020:          2
0x0000000010005024:          2
0x000000001000774c:          2
0x00000000100077a8:          2
0x0000000010007918:          2
0x00000000100098b0:          2
0x00000000100099ac:          2
0x0000000010009a44:          2
0x0000000010009ca8:          2
0x0000000010009f64:          2
0x000000001000a02c:          2
0x00000000100287c4:          2
0x0000000010004f78:          3
0x0000000010007748:          3
0x00000000100077b8:          3
0x00000000100078f0:          3
0x000000001000790c:          3
0x0000000010007910:          3
.
.
0x000000001002876c:          804
0x000000001000963c:         1015
0x0000000010009ebc:         1052
0x0000000010009778:         1168
0x0000000010009434:         1170
0x000000001000977c:         1175
0x0000000010009734:         1204
0x0000000010009644:         1216
0x00000000100287bc:         1420
```

### B.3.4 Frequently Called Routines and Source Code Information

List the frequently called routines and source code information associated with routine.

**Table B–10 Frequently-called Routines**

oct-top /tftpboot/traffic-gen --line		
Function Name:	Total	hits
fdt_check_node_offset:	1	
/opt/tot.testing/executive/libfdt/fdt.c:134:	1	
cvmx_sysinfo_get:	1	
/opt/tot.testing/executive/cvmx-srio.c:405:	1	
fdt_offset_ptr:	1	
/opt/tot.testing/executive/libfdt/fdt.c:81:	1	
cvmx_scratch_read64:	5	
/opt/tot.testing/target/include/cvmx-scratch.h:13:	5	
cvmx_get_cycle:	16	
cvmx_clock_get_count:	16	
/opt/tot.testing/target/include/cvmx-clock.h:112:	16	
cvmx_send_single:	110	
cvmx_write64_uint64:	110	
/opt/tot.testing/target/include/cvmx-access-native.h:476:	110	
fastpath_receive:	281	
/opt/tot.testing/examples/traffic-gen/traffic-gen.c:5864:	67	
/opt/tot.testing/examples/traffic-gen/traffic-gen.c:5866:	38	
/opt/tot.testing/examples/traffic-gen/traffic-gen.c:5931:	37	
/opt/tot.testing/examples/traffic-gen/traffic-gen.c:5877:	30	
/opt/tot.testing/target/include/cvmx-wqe.h:582:	25	
/opt/tot.testing/target/include/cvmx-wqe.h:585:	25	
/opt/tot.testing/examples/traffic-gen/traffic-gen.c:5924:	20	
/opt/tot.testing/examples/traffic-gen/traffic-gen.c:5869:	18	
/opt/tot.testing/examples/traffic-gen/traffic-gen.c:5883:	18	
/opt/tot.testing/examples/traffic-gen/traffic-gen.c:5684:	3	
periodic_update:	380	
/opt/tot.testing/examples/traffic-gen/traffic-gen.c:5320:	150	
/opt/tot.testing/examples/traffic-gen/traffic-gen.c:5341:	150	
/opt/tot.testing/examples/traffic-gen/traffic-gen.c:5322:	64	
.	.	
.	.	
.	.	
packet_transmitter:	18259	
/opt/tot.testing/examples/traffic-gen/traffic-gen.c:6033:	4066	
/opt/tot.testing/examples/traffic-gen/traffic-gen.c:5989:	2361	
/opt/tot.testing/examples/traffic-gen/traffic-gen.c:6070:	2178	
/opt/tot.testing/examples/traffic-gen/traffic-gen.c:6029:	1553	
/opt/tot.testing/examples/traffic-gen/traffic-gen.c:6026:	1531	

### B.3.5 Routines and Hexadecimal PC Values

List the routines and the Hexadecimal PC values associated with routines.

**Table B–11 oct-top: Routines and Hexadecimal PC Values**

\$ oct-top /tftpboot/traffic-gen --address		
Function Name:	Total	Hits
fvwrite_r:	1	
0x000000001014586c:	1	
cvmx_sysinfo_get:	1	
0x000000001003ddf4:	1	
memmove:	1	
0x0000000010138070:	1	
cvmx_scratch_read64:	5	
0x0000000010009730:	5	
cvmx_get_cycle:	14	
cvmx_clock_get_count:	14	
0x000000001001ab38:	14	
cvmx_send_single:	107	
cvmx_write64_uint64:	107	
0x0000000010009758:	107	
fastpath_receive:	289	
0x000000001000791c:	4	
0x0000000010007730:	22	
0x00000000100077c8:	17	
0x00000000100077a0:	14	
0x0000000010007794:	13	
0x00000000100077b0:	13	
0x00000000100077d4:	13	
0x0000000010007764:	11	
0x000000001000775c:	10	
0x0000000010007778:	7	
0x0000000010007788:	7	
0x00000000100077ac:	7	
0x00000000100078e0:	7	
0x0000000010007734:	6	
0x00000000100077d8:	6	
0x0000000010007740:	5	

### B.3.6 Top of Selected Cores

Run top utility on selected cores instead of on all cores running the application.



**Table B-12 oct-top: Running utility on selected Cores**

```
$ oct-top /tftpboot/traffic-gen --core_mask=0x23
Function Name:                Total Hits
cvmx_srio_initialize:         1
kuseg - 32bit User space:    1
cvmx_spinlock_lock:          152
cvmx_spinlock_unlock:        164
process_cmd_packetio:        358
periodic_update:              529
packet_transmitter:           1752
main:                         2882
statistics_gatherer:          2882
fastpath_receive:             5790
```

### B.3.7 OCTEON Hardware Statistics

Display OCTEON hardware statistics that include TADs, DRAM info, FPA pools, and SSO I/O queues.

**Table B-13 OCTEON Hardware Statistics**

```
$ oct-top -s
Core counters:
brmis 63426[0], 32526698[1], 32527062[2], 32528303[3], 32528310[4], 32565565[5],
total (excluding 0) brmis 162675938,
sync 0[0], 0[1], 0[2], 0[3], 0[4], 0[5],
total (excluding 0) sync 0,

L2 statistics for TAD 0:
bus_xmc(addr) count:         22
bus_xmd(store) count:        19
bus_rsc(commit) count:       22
bus_rsd(fill) count:         22
bus_ioc(IO req) count:       22
bus_ior(IO req) count:       22
hit count:                   39650
miss count:                   0
lfb-wait-vab bus utilization: 0%
lfb-wait-lfb bus utilization: 0%

DRAM Statistics:
DRAM ops count:               142357388
dclk count:                   3564342216764
utilization:                   0.00%

FPA Pool statistics:
FPA Pool 0 (Packet Buffers):  816 free
FPA Pool 1 (WQE buffers):     490 free
FPA Pool 2 (PKO cmd buffers):  0 free, PKO output queues 128

POW/SSO Qos Input Queue Statistics
Queue 0: WQE count 0, Buffers allocated 0
Queue 1: WQE count 0, Buffers allocated 0
Queue 2: WQE count 0, Buffers allocated 0
Queue 3: WQE count 0, Buffers allocated 0
Queue 4: WQE count 0, Buffers allocated 0
Queue 5: WQE count 0, Buffers allocated 0
Queue 6: WQE count 0, Buffers allocated 0
Queue 7: WQE count 0, Buffers allocated 0
Total WQEs in all QOS levels: 0 (in-unit & in-memory)
Total Free Buffers in all QOS levels: 243
```

### B.3.8 Cores Running Simple Executive from Linux

Run oct-top on the cores running Simple Executive application booted from Linux.

**Table B-14 oct-top: Simple Executive from Linux**

```
Boot linux on first core
OCTEON ebb6300# tftpboot 0 vmlinux.64
OCTEON ebb6300# bootoctlinux 0 coremask=0x1 mem=512m

$ oct-app-boot namedblock coremask=0x3e
CVMX_SHARED: 0x10190000-0x101a0000
Active coremask = 0x1
INFO: Launching application on coremask 0xff0
Available Coremask = 0x0
$ oct-top named-block -core_mask=0x3e
```

### B.3.9 Frequently Called Routines

List the frequently called routines for Linux (works without passing `vmLinux` file).

**Table B-15 Frequently-called Routines**

\$ oct-top -k		
Function Name:	Total	Hits
__rcu_process_callbacks:	1	
cpumask_next_and:	1	
dequeue_entity:	1	
find_busiest_group:	1	
find_next_bit:	1	
tg_shares_up:	1	
64bit User space:	2719	
r4k_wait:	13589	

### B.3.10 Refresh Rate

Control the refresh rate of the top output.

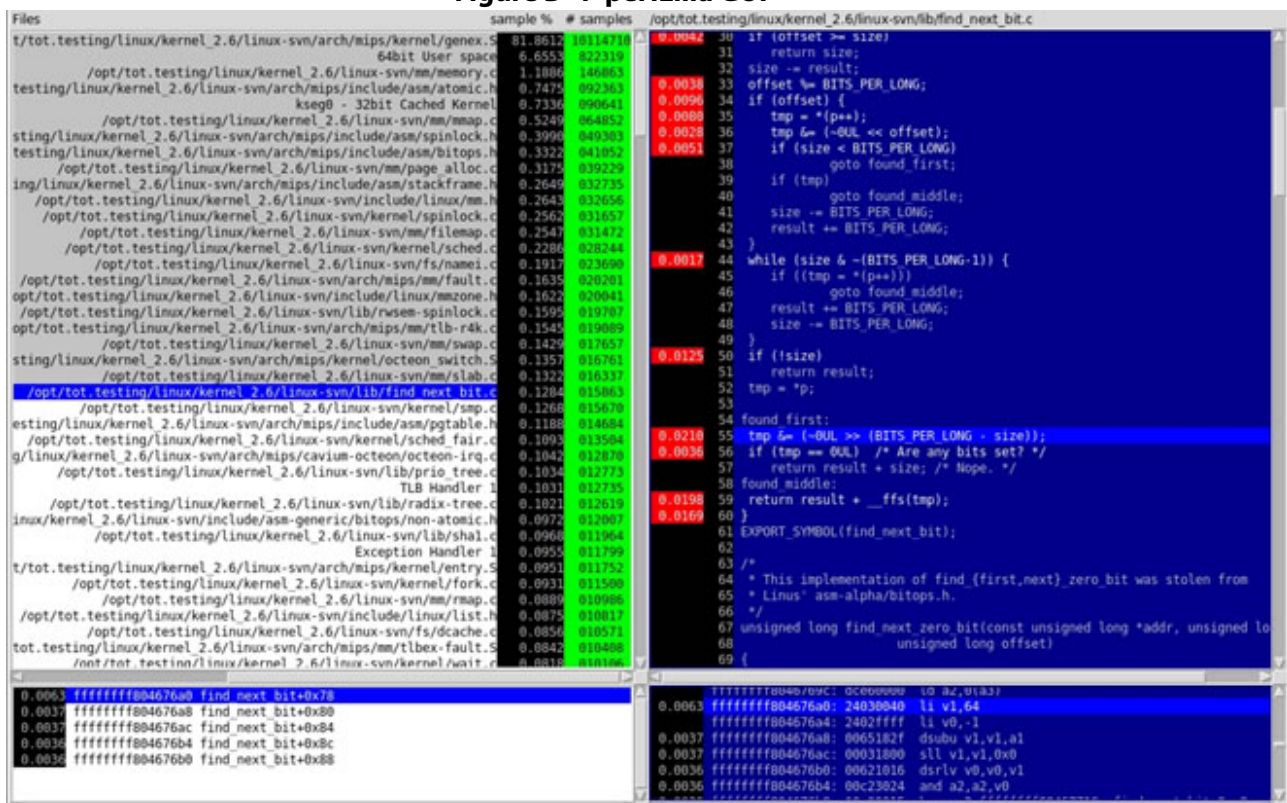
**Table B-16 oct-top: Refresh rate**

```
$ oct-top /tftpboot/-gen --dwell=3
```

## B.4 perfzilla

### B.4.1 Understanding perfzilla output

**Figure B-1 perfzilla GUI**



### B.4.2 Source Code Window

The Source Window display sample %, line number, and the C source code line, as shown in the following table:

**Table B-17 perfzilla: Source Code Window**

0.0210	55	tmp &= (~0UL >> (BITS_PER_LONG - size))
--------	----	---

### B.4.3 Function List window

The Function List window displays the % sample, respective PC value, and function name (including offset), as shown in the following table:

**Table B-18 perfzilla: Function List Window**

0.0063	ffffffff804676a0	find_next_bit+0x78
0.0037	ffffffff804676a8	find_next_bit+0x80
0.0037	ffffffff804676ac	find_next_bit+0x84
0.0036	ffffffff804676b4	find_next_bit+0x8c
0.0036	ffffffff804676b0	find_next_bit+0x88

### B.4.4 ASM window

The ASM window displays information about the disassembly of the selected C source code, as shown in the following table:

**Table B-19 perfzilla: ASM Window**

0.0063	ffffffff804676a0:	24030040	livl,64
	ffffffff804676a4:	2402ffff	liv0,-1
0.0037	ffffffff804676a8:	0065182f	dsubuvl,v1,a1
0.0037	ffffffff804676ac:	00031800	sllvl,v1,0x0
0.0036	ffffffff804676b0:	00621016	dsrlvv0,v0,v1
0.0036	ffffffff804676b4:	00c23024	anda2,a2,v0
0.0036	ffffffff804676b8:	10c00015	beqza2,ffffffff80467710
<find_next_bit+0xe8>			

The ASM window updates information about the disassembly of the selected C source code in the source file, as shown in the following table:

**Table B-20 perfzilla: ASM Window - Source Code Disassembly**

tmp &= (~0UL >> (BITS_PER_LONG - size))	
li v1,64	
store the value of BITS_PER_LONG in v1 register	
li v0,-1	
store ~0UL value in to v0 register	
dsubu v1,v1,a1	
v1 = v1 - a1 (v1 is 64 and a1 refers to the subroutine find_next_bit argument size)	
sll v1,v1,0x0	
dsrlv v0,v0,v1	
v0 = ~0uL >> (BITS_PER_LONG - size)	
and a2,a2,v0	
a2 = a2 & v0	